

SMARTeBuses

Storing Data in a Hadoop Cluster

Deliverable number: WP2-D1

Project Acronym: SMARTeBuses
Project Full Title: SMART electric Buses
Grant Number: 19/RDD/519
Project URL: <https://smartebuses.github.io/web/>

Deliverable type:	Deliverable
Dissemination level:	Public (PU)
Delivery Date:	August 31st, 2020
Number of pages:	65
Keywords:	Big Data, Distributed File System, Large-scale Data Analysis, SMARTe-Buses, Artificial Intelligence
Authors:	Dr. Ignacio Castiñeiras, Cork Institute of Technology
Peer review:	Dr. Alejandro Arbelaez, University College Cork

Contents

1	Introduction	5
2	Big Data Ecosystem	6
2.1	Hadoop	6
2.1.1	Hadoop Distributed File System	6
2.1.2	Hadoop Yarn	8
2.1.3	Hadoop MapReduce	8
2.2	Spark	12
2.2.1	Data Storage and Resource Manager	12
2.2.2	Spark Core	13
2.2.3	Spark SQL	14
3	Configuration and Code Examples	16
3.1	Local Cluster	16
3.2	Java	16
3.3	Python	17
3.4	Hadoop	19
3.4.1	Installing Hadoop	19
3.4.2	Start a Hadoop Cluster	22
3.4.3	Dataset	23
3.4.4	HDFS	24
3.4.5	MapReduce	31
3.4.6	Stop a Hadoop Cluster	40
3.5	Spark	41
3.5.1	Installing Spark	41
3.5.2	Spark Core	42
3.5.3	Spark SQL	51
4	Conclusions and Future Work	62

List of Figures

2.1	HDFS and Local File System	7
2.2	File to Block and Block Replication	7
2.3	MapReduce Application Phases	10
2.4	Map Phase	11
2.5	Reduce Phase	11
2.6	Spark Components	12
2.7	Spark Components	13
2.8	Spark Core: User Program	14
2.9	Spark Core: Program Execution	14
3.1	Script openJDK_8.sh	16
3.2	Script python_3_7_7.sh	18
3.3	Script Hadoop_2_7_1.sh	20
3.4	Script Start_Hadoop_Cluster.sh	22
3.5	Script data_analysis.sh	24
3.6	Cluster Overview	25
3.7	HDFS	26
3.8	Dataset Folder: my_dataset	27
3.9	Command put: my_dataset from local file system to HDFS	28
3.10	HDFS With New Folder my_result	29
3.11	Command get: my_result from HDFS to local file system	30
3.12	ResourceManager Overview	31
3.13	ResourceManager: MapReduce Application in Progress	32
3.14	ResourceManager: MapReduce Application Finished	33
3.15	HDFS: MapReduce Result in my_result	34
3.16	HDFS: MapReduce Result Brought Back to Local File System	35
3.17	my_mapper.py: Import Section	36
3.18	my_mapper.py: my_map Function	37
3.19	my_mapper.py: Main Entry Point	38
3.20	my_mapper.py: Import Section	38
3.21	my_mapper.py: my_map Function	39
3.22	my_mapper.py: Main Entry Point	40
3.23	Script Stop_Hadoop_Cluster.sh	40
3.24	Script Spark_2_4_5.sh	41
3.25	my_Spark_Core_example.py: Import Section	43
3.26	my_Spark_Core_example.py: my_spark_core_job Function	44
3.27	my_Spark_Core_example.py: Main Entry Point	45
3.28	inputRDD Content	46
3.29	all_wordsRDD Content	46
3.30	clean_wordsRDD Content	46
3.31	RDDs Content	46

3.32	ResourceManager: Spark Core Application in Progress	47
3.33	ResourceManager: Spark Core Application Finished	48
3.34	HDFS: Spark Core Result in my_result	49
3.35	HDFS: Spark Core Result Brought Back to Local File System	50
3.36	my_Spark_SQL_example.py: Import Section	51
3.37	my_Spark_SQL_example.py: my_spark_sql_job Function	53
3.38	my_Spark_SQL_example.py: my_spark_sql_job Function	54
3.39	my_Spark_SQL_example.py: Main Entry Point	55
3.40	inputDF Content	56
3.41	sentenceDF Content	56
3.42	wordsDF Content	56
3.43	cleanDF Content	57
3.44	lowerDF Content	57
3.45	solutionDF Content	57
3.46	ResourceManager: Spark SQL Application in Progress	58
3.47	ResourceManager: Spark SQL Application Finished	59
3.48	HDFS: Spark SQL Result in my_result	60
3.49	HDFS: Spark SQL Result Brought Back to Local File System	61

List of Acronyms

HDFS Hadoop Distributed File System

JRE Java Runtime Environment

Yarn Yet Another Resource Negotiator

DSL Domain Specific Language

RDD Resilient Distributed Datasets

1 Introduction

Big Data is the ability of society to harness information in novel ways to produce useful insights or goods and services of significant value [1]. In this work package we focus on the use of Big Data technology to store, analyse and collect online public data related to the bus network (e.g., routes, timetables, delays) and wind power available from wind farms in the country. Furthermore, we focus on developing a realistic reference road network with essential information for a strategic planning and design of an electric bus network in Ireland. In line with the goals of the work package, we introduce the Big Data ecosystem of tools we are going to use for storing and analysis bus and wind power-related large-datasets.

The document is structured as follows:

Chapter 2 presents our ecosystem of tools to store and analyse large-scale datasets. First, Section 2.1 presents Apache Hadoop [2], specifically its components Hadoop Distributed File System (HDFS) [3, 4], Hadoop Yarn [5, 6] and Hadoop MapReduce [7, 8] for storing and analysing datasets. Then, Section 2.2 presents Apache Spark [9] as our primary data analytics tool (relying on HDFS and Yarn for data storage and job management, resp). In particular, it presents its components Spark Core [10, 11] and Spark SQL [12, 13]. While Python is selected as the programming language of choice (with both MapReduce and Spark providing an API for it), the data analytics applications will run on top of the Java Runtime Environment (JRE) [14].

Chapter 3 serves as a tutorial for installing and configuring such ecosystem of tools, as well as presenting some introductory data analysis examples. First, Section 3.1 presents the specs of the local server. Then, sections 3.2, 3.3, 3.4 and 3.5 present the setup of Java, Python, Hadoop and Spark, resp., including some introductory data analysis examples running on the local cluster.

2 Big Data Ecosystem

This chapter presents our ecosystem of tools to store and analyse large-scale datasets.

2.1 Hadoop

Apache Hadoop is a data framework for store and process big data distributed on clusters of commodity machines. Using the definition of its own website: *“The Apache Hadoop software library is a framework that allows for the distributed processing of large data sets across clusters of computers using simple programming models. It is designed to scale up from single servers to thousands of machines, each offering local computation and storage. Rather than rely on hardware to deliver high-availability, the library itself is designed to detect and handle failures at the application layer, so delivering a highly-available service on top of a cluster of computers, each of which may be prone to failures”* [2].

The main idea of the framework is for data to be distributed among the nodes of the cluster for both its storage and processing. This way, each individual node works as much as possible on the data it hosts, minimising the need to talk to other nodes of the cluster (and thus minimising the data transferred over the network). Data is replicated among different nodes of the cluster for redundancy, availability and fault tolerant-free computations. If a node is unexpectedly shut down, the system is still ready to store and process the entire dataset without suffering a noticeable penalisation in the performance to achieve it. At failure, the workload of the node is assumed by any other alive node of the cluster, without any loss of data. Thus, the outcome of the computation does not become affected. When the node recovers, it just joins again the set of operative nodes, becoming available again for storing and processing data. Likewise, the addition of new nodes to the cluster is easy, and provides extra storage and processing capacity under demand. As the nodes cooperate on the tasks, the system is capable of dealing with peak workloads, which result only in slightly penalisation for the overall performance.

There are 3 components of Hadoop that we are going to use:

- HDFS: A distributed file system designed to efficiently allocate data across the multiple commodity machines (nodes) of the cluster.
- Yet Another Resource Negotiator (Yarn): A resource manager, responsible for schedule and monitor the execution of our data analysis applications.
- MapReduce: A framework for easily writing applications processing large-scale datasets across a cluster in a reliable, fault-tolerant manner.

2.1.1 Hadoop Distributed File System

We use HDFS as a distributed file system to store our datasets across a cluster. The master/slave architecture of HDFS consists of two types of daemons:

- A single NameNode. This is a master daemon managing the file system namespace and regulating access to files. In brief, the NameNode acts as the table of contents: it contains no single file itself, but it knows all files in the filesystem, as well as their distribution among the nodes.

- A number of DataNodes, usually one per node in the cluster. This is a slave daemon managing the storage of the node it runs on. That is, the DataNode controls all the files stored in the node, but it lacks any knowledge about the data stored in other nodes.

Figure 2.1 presents a cluster with 4 nodes, where 1 node contains the NameNode and the other 3 a DataNode each. As we can see, the memory storage of each node contains its own local file system and its fraction/slice of the HDFS.

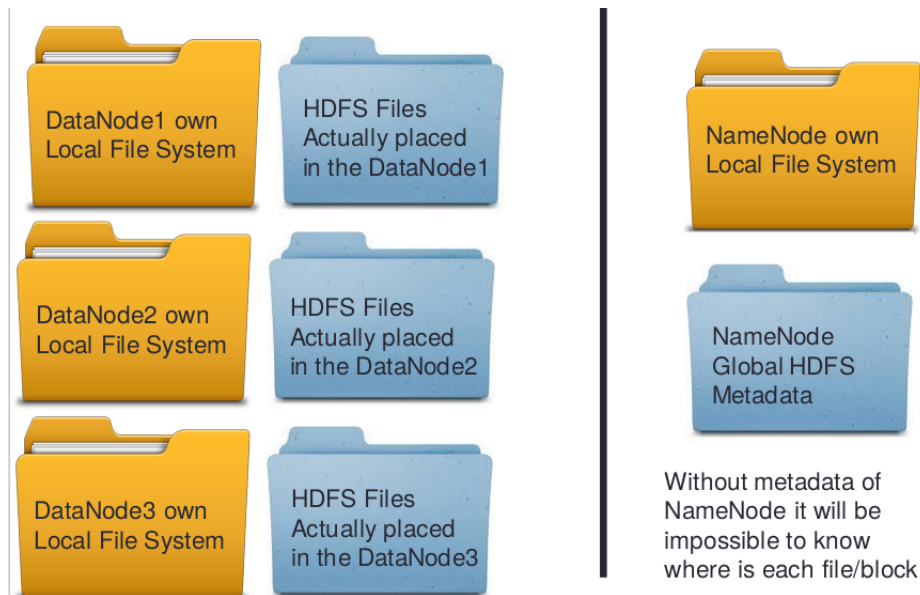


Figure 2.1: HDFS and Local File System

HDFS is designed to reliably store very large files. Its approach consists on splitting each file into manageable blocks and replicate each blocks among multiple nodes (typically 3, although this is configurable). Figure 2.2 shows a file split into 3 blocks when being written to HDFS and the replication of each block among the DataNodes.




Figure 2.2: File to Block and Block Replication

We envision our data analysis applications to follow a very simple flow:

- Read the dataset from a logic folder of HDFS (e.g., `my_dataset`).

- Write the data analysis result to a new logic folder of HDFS (e.g., `my_result`).

The above requires us to:

- Create a new logic folder `my_dataset` into HDFS.
- Bring `my_dataset` from the local file system to the HDFS folder.
- Bring `my_result` from HDFS to the local file system.

HDFS provide the specific commands `mkdir`, `put` and `get` to accomplish the aforementioned tasks, resp. In brief, any read/write operation is coordinated by the NameNode, who refers to the associated DataNode of the node hosting the block for managing it.

2.1.2 Hadoop Yarn

We use the Yarn system as a resource manager, responsible for schedule and monitor the execution of our data analysis applications.

The key idea of Yarn is to split up the functionalities of resource management and job scheduling/monitoring into separate daemons:

- A single ResourceManager. This is a master daemon distributing the resources of the cluster among all the applications. In an analogy with HDFS, it will play the role of the NameNode, in the sense it executes nothing by itself, but it knows all the resources and tasks to be executed among the nodes of the cluster.
- A NodeManager per node in the cluster. This is a slave daemon monitoring the resource usage (cpu, memory, disk, network) of the node it runs on. In an analogy with HDFS, it will play the role of the DataNode, in the sense it controls the resources of the node and monitors the execution of the tasks assigned to it, but it lacks any knowledge about the resources and execution in other nodes of the cluster.

In brief, given a MapReduce or Spark application, the ResourceManager will provide the application with resources from 1 or more nodes of the cluster. Then, each NodeManager involved in the execution of the application will accomplish its assigned tasks with the resources provided.

2.1.3 Hadoop MapReduce

We use MapReduce as our first framework for easily writing applications processing large-scale datasets across a cluster in a reliable, fault-tolerant manner. A MapReduce application can be seen as a pipeline process operating on files that uses streaming for communication and high-level programming features to isolate the data processing from both the cluster and the pipeline complexities.

As previously stated, we envision a MapReduce application to read the dataset from a HDFS folder (e.g., `my_dataset`) and produce its results to a novel HDFS folder (e.g., `my_result`). In doing so, the application is to follow 3 phases: `map`, `sort` and `reduce`. Figure 2.3 presents an introductory word count MapReduce application, viewing its different phases from a logic and a physical point of views, resp.

As we can see, the *Map Phase* processes the entire dataset of `my_dataset` in a completely parallel manner. On doing so, each data block is processed on its own node as much as possible, with typically a `Map` process allocated to each block. The `Map` functionality is to be programmed just once, with all the processes applying it to their respective data slice (cf. middle picture of Figure 2.3). While the right level of parallelism for `Map` processes seems to be around 10-100 per node, in general the framework figures out the number of processes to be applied, as well as the block allocation for them.

This abstracts the cluster complexity from the user, who just has to focus on the programming of the Map functionality.

To specify the Map functionality we make use of Hadoop Streaming [15], which abstracts the Map process from its underlying programming language by only requiring it to be in the form of an executable command. In our case, we choose to program the Map functionality in the form of a Python script (e.g., `my_mapper.py`). Figure 2.4 shows `my_mapper.py` as a black-box. As we can see, it must receive its input from the standard input, and produce a bunch of key-value pairs by writing them to the standard output.

Coming back to Figure 2.3, we see that the *Sort Phase* sorts the outputs of the Map processes, which are then input to the Reduce processes. Again, while the right level of parallelism for Reduce processes seems to be around 1-2 per node, this is abstracted to the user who just has to focus on the programming of the Reduce functionality. What it is ensured is that all sorted entries having the same key are to be assigned to the same Reduce process.

As we did with the Map functionality, we use Hadoop Streaming to program the Reduce functionality in the form of a Python script (e.g., `my_reducer.py`). Figure 2.5 shows `my_reducer.py` as a black-box. As we can see, it must receive the key-value pairs produced by Map as input from the standard input, and produce a bunch of key-value pairs by writing them to the standard output.

Finally, the bulk of results from the *Reduce Phase* produce the new HDFS folder `my_result`. In particular, the output of each Reduce process turns into one or more new blocks in the folder.

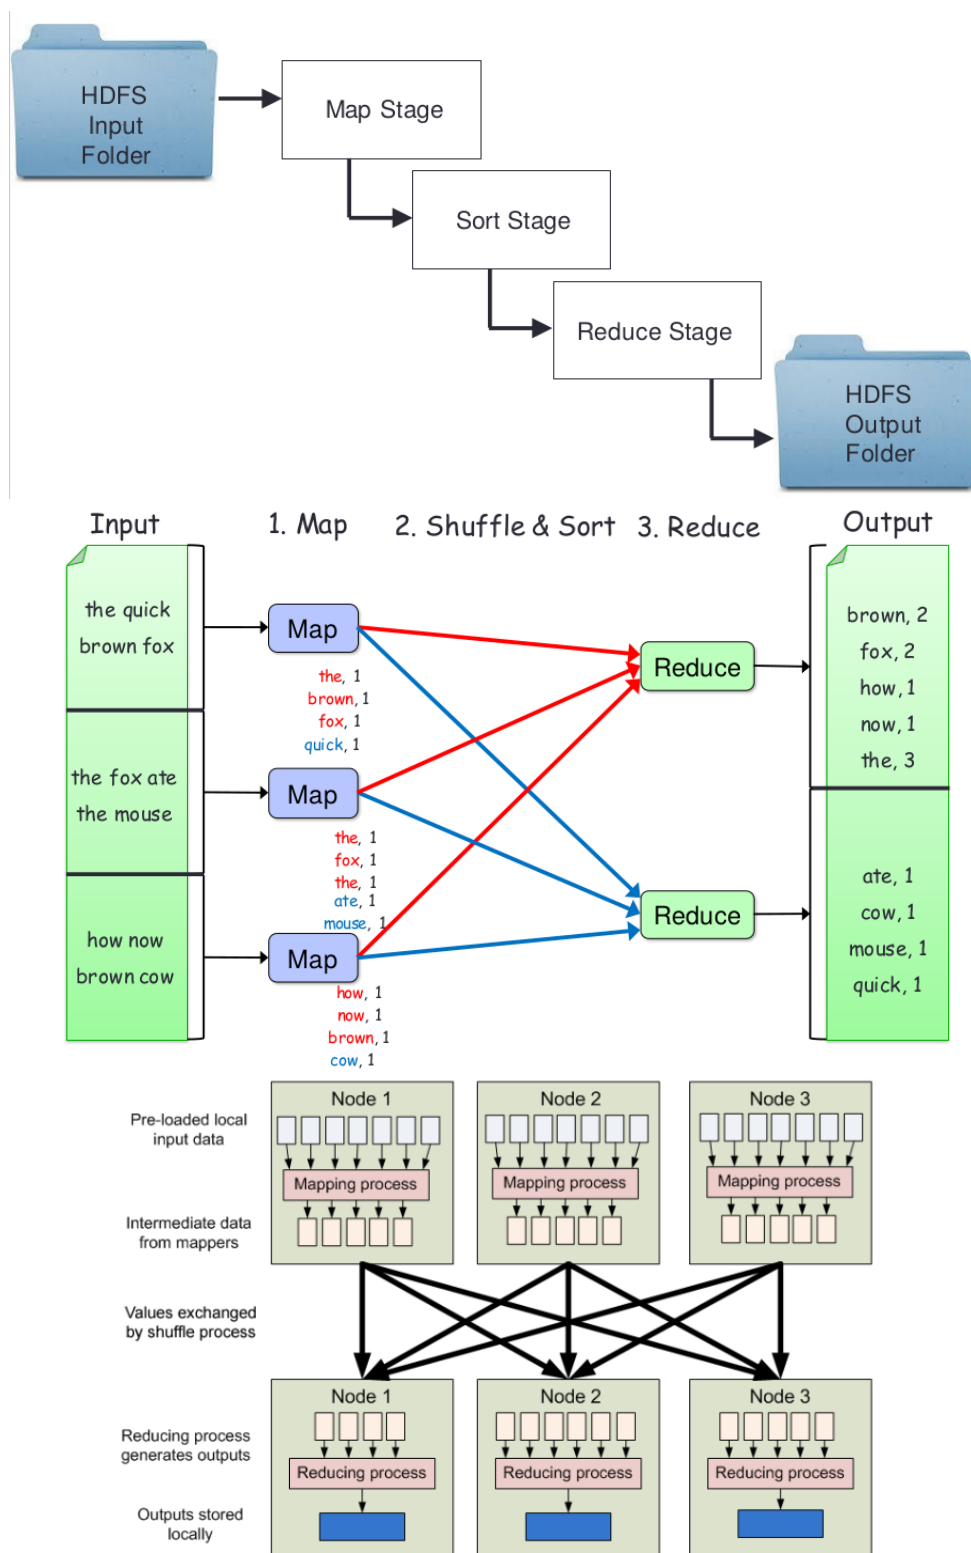


Figure 2.3: MapReduce Application Phases




Figure 2.4: Map Phase




Figure 2.5: Reduce Phase

2.2 Spark

We use Apache Spark as our second framework for easily writing applications processing large-scale datasets across a cluster in a reliable, fault-tolerant manner. Using the definition of its own website: *Apache Spark is a unified analytics engine for large-scale data processing* [9]. It is an open-source, distributed, general-purpose, cluster-computing framework designed for 3 purposes:

- Be easy to use, allowing us to develop applications locally, using a high-level API.
- Be fast, enabling interactive use and complex algorithms.
- Be general, allowing us to combine multiple types of computations, including text, SQL, graph and machine learning processing (both offline and online) that might previously have required different engines.

Spark itself is written in Scala [16], and runs on the Java Virtual Machine (JVM). However, it offers simple APIs in Scala, Java, Python and R. In our case, we choose Python to program our data analysis applications.

Figure 2.6 presents the different components of Spark. We are going to rely on HDFS for data storage and on Yarn for resource management and job scheduling. For developing our data analytics applications we are going to use Spark Core and Spark SQL.

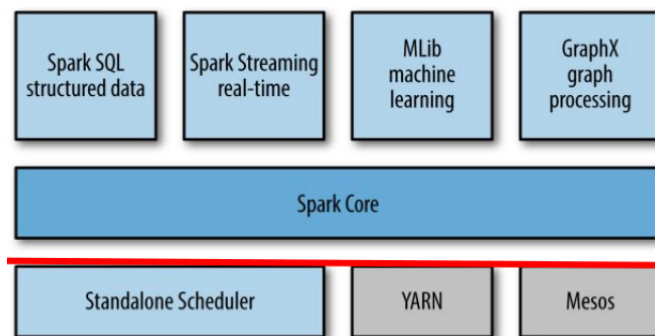


Figure 2.6: Spark Components

2.2.1 Data Storage and Resource Manager

On its own, Spark is not a data storage solution; it performs computations on Java Virtual Machines (JVMs) that last only for the duration of a Spark application. While Spark can be run locally on a single machine with a single JVM (called local mode), this mode is only useful for debugging and testing purposes. Once the application is tested, Spark is designed to efficiently scale up from one to many thousands of compute nodes, so we will use it in distributed mode across a cluster.

When used in a cluster, Spark is used in tandem with a distributed storage system (in our case HDFS) and with a cluster manager (in our case Yarn).

- HDFS is used to provide the input dataset used by Spark application (e.g., the HDFS folder `my_dataset`), as well as to store the results such application produces (e.g., the HDFS folder `my_result`).
- Yarn is used to schedule the execution of a Spark application, and for assigning and monitoring the resources of each node executing it.




Figure 2.7: Spark Components

Figure 2.7 presents a general view of the execution of a Spark application in a cluster. The master/slave architecture of Spark consists of two types of daemons:

- A single SparkDriver. This is a master daemon coordinating the execution of the User program. In an analogy to HDFS, it plays the role of the NameNode, in the sense that it does not compute anything by itself, but controls who is executing each task and how is this execution going.
- A number of SparkExecutors. This is a slave daemon managing the execution of a single given task. A number of SparkExecutors are distributed among the cluster. While the number of cores per executor can be configured at the user program, typically they correspond to the physical cores on a node, and an executor cannot span cores of different nodes. In an analogy to HDFS, a SparkExecutor plays the role of the DataNode, in the sense that it carries out the computation of the task being given (and reports its status to the SparkDriver), but lacks any knowledge about the tasks executed by other SparkExecutors.

2.2.2 Spark Core

Spark Core contains the basic functionality of Spark, including:

- The main data abstraction being offered to users to express their programs: The Resilient Distributed Datasets (RDD)
- The implementation of the SparkDriver and SparkExecutor daemons needed to execute a Spark application. This includes the Spark driver functionality to generate a logical plan (Direct Acyclic Graph per job, with concrete stages and tasks) and a physical plan (schedule and tracking of the tasks execution).
- Other functionality such as memory management, fault recovery and the interaction with storage systems.

An RDD simply defines a collection of items. It is:

- Indivisible (logically presented as an atomic variable).
- Generic, but statically-typed (available for any data type T as long as the type T sticks for all its elements).
- Lazily-evaluated (only computed if required, and as much as required).
- Non-mutable (cannot change type nor value).

In this context, we can think of an RDD as a list (in the sense that its elements can be repeated) or a set (in the sense that its elements have no particular default order).

The API of Spark Core offers Creation, Transformation, Persistent and Observer operations. Figure 2.8 presents the life-cycle of a Spark Core user program, based on these operations.

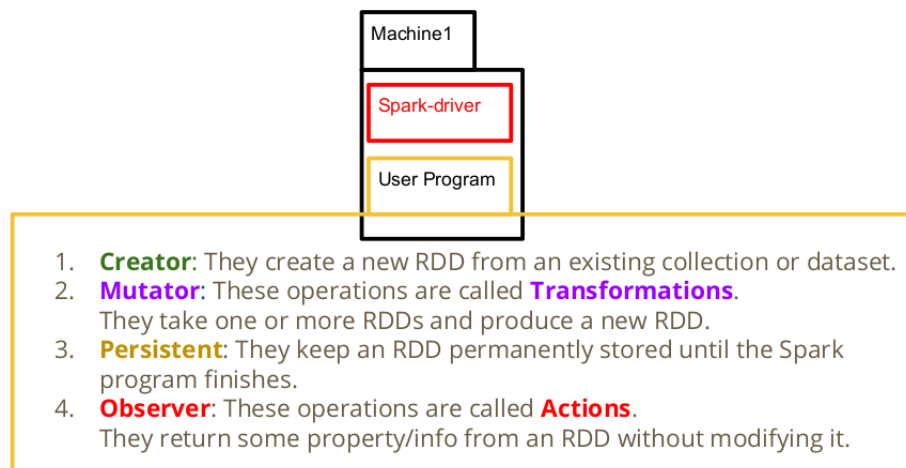


Figure 2.8: Spark Core: User Program

Figure 2.9 presents the execution of a Spark Core application. The application is defined as a set of Jobs triggered by the action operations of the user program. Each Job consists on a sequential execution of stages. Each new stage is caused by a shuffle of information among nodes executing different tasks. While no external communication is needed, each SparkExecutor works locally, performing a pipeline of tasks.

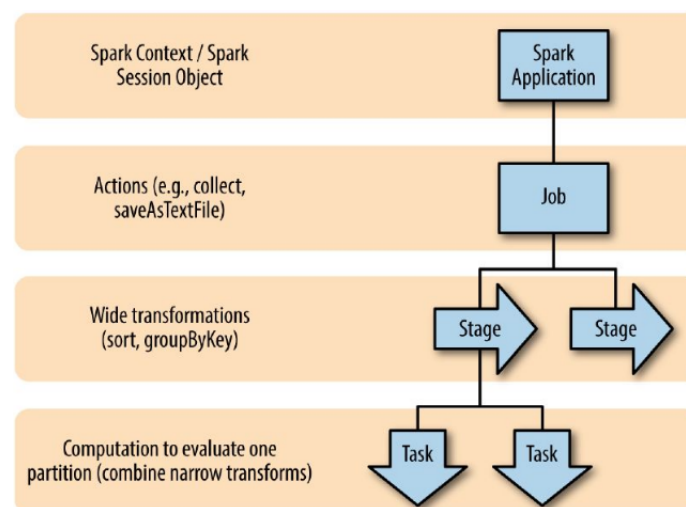


Figure 2.9: Spark Core: Program Execution

2.2.3 Spark SQL

Spark SQL is the module integrating relational processing with the functional programming API of Spark. By using it, we can benefit of a higher-level data abstraction (DataFrames) for ingesting, querying and persisting (semi)structured data using relational queries via a Domain Specific Language (DSL). Under the hood, Spark SQL translates a DataFrame-based program into an equivalent RDD-based one via a catalyst optimiser (which semantically analyse the query expressed by the user) and a Tungsten encoder optimiser (which translates the query to an equivalent binary RDD-based format). An DataFrame simply defines a collection of items. It is:

- Indivisible (logically presented as an atomic variable).
- Structured, in the sense of having a fixed number of fields, each of them of a concrete data type T.
- Lazily-evaluated (only computed if required, and as much as required).
- Non-mutable (cannot change type nor value).

In this context, we can think of a DataFrame as a table in a relational database (in the sense that each Row follows the schema) or a collection in a NoSQL document oriented database (in the sense that the content of the collection is distributed).

Similarly to Spark Core, the API of Spark SQL offers Creation, Transformation, Persistent and Observer operations. Spark SQL provides them via a catalog of DSL operators. This DSL is being updated on each new release of Spark, making Spark SQL more and more expressive, and thus making it easier for the user to develop its data analysis applications.

3 Configuration and Code Examples

This chapter serves as a tutorial for installing and configuring the ecosystem of tools, as well as presenting some introductory data analysis examples.

3.1 Local Cluster

The local cluster used in this project consist on a server with the following specs:

- **Processor:** Intel Xeon W-2175 2.5GHz, 4.3GHz Turbo, 14C, 19.25M Cache, HT, (140W) DDR4-2666.
- **RAM:** 64GB 4x16GB DDR4 2666MHz RDIMM ECC Memory.
- **Hard Drive:** M.2 2TB PCIe NVMe Class 40 Solid State Drive.
- **Operating System:** Ubuntu 20.04 LTS [17].

The server has been purchased with the funding for equipment assigned to the project.

3.2 Java

In this section we discuss how to install and configure Java OpenJDK 8 [18]. While there is a more recent version available (Java OpenJDK 11 [19]), Java OpenJDK 8 is the most recent version supported by the latest stable release of Spark, thus making it our JRE of choice.

Figure 3.1 presents the script to install and configure Java OpenJDK 8. It can be run from a terminal in our local server.

```
(01) #!/bin/bash
(02) sudo apt-get update
(03) sudo apt-get install openjdk-8-jdk
(04) sudo update-alternatives --config java
(05) sudo update-alternatives --config javac
(06) sudo gedit ~/.bashrc
```

Figure 3.1: Script openJDK_8.sh

Next, we present more detailed instructions about the steps followed in the script:

- (01) #!/bin/bash
We indicate bin bash as the interpreter being used.
- (02) sudo apt-get update
We update apt-get, the tool to handle packages in Linux.

- (03) `sudo apt-get install openjdk-8-jdk`

We install Java OpenJDK 8.

- (04) `sudo update-alternatives --config java`

We ensure that Java OpenJDK 8 is the default JRE. The command shows the different JREs installed in the system, highlighting with a symbol `*` the default one. For example, in the case below both Java OpenJDK 8 and Java OpenJDK 11 are installed in the system, with Java OpenJDK 8 being the default one being used.

Selection	Path

0	/usr/lib/jvm/java-11-openjdk-amd64/bin/java
1	/usr/lib/jvm/java-11-openjdk-amd64/bin/java
* 2	/usr/lib/jvm/java-8-openjdk-amd64/jre/bin/java

- (05) `sudo update-alternatives --config javac`

We ensure that Java OpenJDK 8 is the default Java compiler. The command shows the different Java compilers installed in the system, highlighting with a symbol `*` the default one. For example, in the case below both Java OpenJDK 8 and Java OpenJDK 11 are installed in the system, with Java OpenJDK 8 being the default one being used.

Selection	Path

0	/usr/lib/jvm/java-11-openjdk-amd64/bin/javac
1	/usr/lib/jvm/java-11-openjdk-amd64/bin/javac
* 2	/usr/lib/jvm/java-8-openjdk-amd64/bin/javac

- (06) `sudo gedit ~/.bashrc`

We use super user privileges with the text editor `gedit` so as to modify the content of the configuration file `bashrc`. In our case we add:

```
export JAVA_HOME="/usr/lib/jvm/java-1.8.0-openjdk-amd64"
export PATH=$PATH:/usr/lib/jvm/java-1.8.0-openjdk-amd64/bin
```

3.3 Python

In this section we discuss how to install and configure Python 3.7.7 [20]. While there is a more recent version available (Python 3.8.2 [21]), Python 3.7.7 is the most recent version supported by the latest stable release of Spark, thus making it our programming language of choice.

Figure 3.2 presents the script to install and configure Python 3.7.7. It can be run from a terminal in our local server.

Next, we present more detailed instructions about the steps followed in the script:

- (01) `#!/bin/bash`

We indicate bin bash as the interpreter being used.

- (02) `sudo apt update`

We update apt, the tool for managing deb packages in Ubuntu.

- (03) `sudo apt install build-essential zlib1g-dev libncurses5-dev \`
`libgdbm-dev libnss3-dev \`

```

(01) #!/bin/bash
(02) sudo apt update
(03) sudo apt install build-essential zlib1g-dev \
        libncurses5-dev libgdbm-dev libnss3-dev \
        libssl-dev libreadline-dev \
        libffi-dev libsqlite3-dev \
        wget libbz2-dev
(04) wget https://www.python.org/ftp/python/3.7.7/Python-3.7.7.tgz
(05) tar -xf Python-3.7.7.tgz
(06) cd Python-3.7.7
(07) ./configure --enable-optimizations
(08) make -j $(nproc)
(09) sudo make altinstall
(10) sudo python3.7 -m pip install -U pip
(11) which python3.7
(12) python3.7

```

Figure 3.2: Script python_3.7.7.sh

```

        libssl-dev libreadline-dev \
        libffi-dev libsqlite3-dev \

        wget libbz2-dev

```

We use apt to install the required dependency packages.

- (04) `wget https://www.python.org/ftp/python/3.7.7/Python-3.7.7.tgz`

We download Python 3.7.7 as a Gzipped source tarball.

- (05) `tar -xf Python-3.7.7.tgz`

We extract it.

- (06) `cd Python-3.7.7`

We move to the extracted Python 3.7.7 folder.

- (07) `./configure --enable-optimizations`

We run the configure script, which looks for dependencies.

- (08) `make -j $(nproc)`

We run the make across the number of processors we have in the server.

- (09) `sudo make altinstall`

We use `altinstall` instead of `install` so as to allow Python 3.7.7 to co-exist with other potential versions installed in the server.

- (10) `sudo python3.7 -m pip install -U pip`

We upgrade pip, the package installer for Python, to its most recent version 20.1 [22].

- (11) `which python3.7`

We get the location of Python3.7.7 in our server. In this case it is:

`/usr/local/bin/python3.7`

- (12) `python3.7`

We launch Python3.7.7 to ensure it has been installed. In this case we should get the following:

```
Python 3.7.7 (default, May 6 2020, 16:41:07)
[GCC 7.5.0] on linux
Type "help", "copyright", "credits" or "license"
for more information.
>>>
```

3.4 Hadoop

In this section we discuss how to install, configure and manage a Hadoop cluster. We also present a MapReduce application running on the cluster.

3.4.1 Installing Hadoop

We install and configure Hadoop 2.7.1 [23]. While there is a more recent version available (Hadoop 3.2.1 [24]), Hadoop 2.7.x is the most recent version supported by the latest stable release of Spark, thus making it our version of choice.

Figure 3.3 presents the script to install and configure Hadoop 2.7.1 as a Single Node Cluster with Pseudo-Distributed Operation. We choose this mode as our cluster contains just 1 node. Thus, all the HDFS, Yarn, MapReduce and Spark daemons presented in Chapter 2 still take place, each of them running as an independent process. The script can be run from a terminal in our local server.

Next, we present more detailed instructions about the steps followed in the script:

- (01) `#!/bin/bash`

We indicate bin bash as the interpreter being used.

- (02) `sudo apt-get install openssh-server openssh-client`

We need to be able to access localhost by ssh without a passphrase. Thus, first we use apt-get to install OpenSSH, which is the premier connectivity tool for remote login with the ssh protocol [25].

- (03) `ssh-keygen -t rsa -P ""`

We use ssh-keygen to generate a public/private rsa key pair. If run successfully, it saves the identification and public key in `$HOME/.ssh/`.

- (04) `cat $HOME/.ssh/id_rsa.pub >> $HOME/.ssh/authorized_keys`

We use the command cat to append the public keys (generated by ssh-keygen) to our file of listed authorised keys.

- (05) `ssh localhost`

We ensure we can connect now by ssh to the localhost without a passphrase.

- (06) `wget https://archive.apache.org/dist/hadoop/common/hadoop-2.7.1/hadoop-2.7.1.tar.gz`

We download Hadoop 2.7.1 as a Gzipped source tarball.

- (07) `tar -xzf hadoop-2.7.1.tar.gz`

We extract it.

```

(01) #!/bin/bash
(02) sudo apt-get install openssh-server openssh-client
(03) ssh-keygen -t rsa -P ""
(04) cat $HOME/.ssh/id_rsa.pub >> $HOME/.ssh/authorized_keys
(05) ssh localhost
(06) wget https://archive.apache.org/dist/hadoop/
      common/hadoop-2.7.1/hadoop-2.7.1.tar.gz
(07) tar -xzvf hadoop-2.7.1.tar.gz
(08) sudo mv hadoop-2.7.1 /usr/local/hadoop/
(09) sudo gedit ~/.bashrc
(10) sudo gedit /usr/local/hadoop/etc/hadoop/hadoop-env.sh
(11) sudo gedit /usr/local/hadoop/etc/hadoop/core-site.xml
(12) sudo gedit /usr/local/hadoop/etc/hadoop/hdfs-site.xml
(13) sudo gedit
      /usr/local/hadoop/etc/hadoop/mapred-site.xml.template
(14) sudo cp
      /usr/local/hadoop/etc/hadoop/mapred-site.xml.template
      /usr/local/hadoop/etc/hadoop/mapred-site.xml
(15) sudo gedit /usr/local/hadoop/etc/hadoop/yarn-site.xml
(16) hadoop

```

Figure 3.3: Script Hadoop_2.7.1.sh

- (08) `sudo mv hadoop-2.7.1 /usr/local/hadoop/`

We move the extracted folder to `/usr/local/hadoop`.

- (09) `sudo gedit ~/.bashrc`

We use super user privileges with the text editor `gedit` so as to modify the content of the configuration file `bashrc`. In our case we add:

```

export PATH=$PATH:/usr/local/hadoop/bin:/usr/local/hadoop/sbin/
export HADOOP_HOME=/usr/local/hadoop/
export HADOOP_CONF_DIR=/usr/local/hadoop/etc/hadoop
export HADOOP_MAPRED_HOME=/usr/local/hadoop/
export HADOOP_COMMON_HOME=/usr/local/hadoop/
export HADOOP_HDFS_HOME=/usr/local/hadoop/
export YARN_HOME=/usr/local/hadoop/
export HADOOP_COMMON_LIB_NATIVE_DIR=/usr/local/hadoop/lib/native
export HADOOP_OPTS="-Djava.library.path=/usr/local/hadoop/lib"
export JAVA_LIBRARY_PATH=$HADOOP_HOME/lib/native:$JAVA_LIBRARY_PATH

```

- (10) `sudo gedit /usr/local/hadoop/etc/hadoop/hadoop-env.sh`

We use super user privileges with the text editor `gedit` so as to modify the content of the configuration file `hadoop-env.sh`. In our case we add:

```

export JAVA_HOME=/usr/lib/jvm/java-1.8.0-openjdk-amd64

```

- (11) `sudo gedit /usr/local/hadoop/etc/hadoop/core-site.xml`

We use super user privileges with the text editor `gedit` so as to modify the content of the configuration file `core-site.xml`. In our case we fill the (originally empty `<configuration></configuration>`) configuration section with:

```
<configuration>
<property>
<name>fs.defaultFS</name>
<value>hdfs://localhost:9000</value>
</property>
</configuration>
```

- (12) `sudo gedit /usr/local/hadoop/etc/hadoop/hdfs-site.xml`

We use super user privileges with the text editor `gedit` so as to modify the content of the configuration file `hdfs-site.xml`. In our case we fill the (originally empty `<configuration></configuration>`) configuration section with:

```
<configuration>
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
</configuration>
```

- (13) `sudo gedit /usr/local/hadoop/etc/hadoop/mapred-site.xml.template`

We use super user privileges with the text editor `gedit` so as to modify the content of the configuration file `mapred-site.xml.template`. In our case we fill the (originally empty `<configuration></configuration>`) configuration section with:

```
<configuration>
<property>
<name>mapreduce.framework.name</name>
<value>yarn</value>
</property>
</configuration>
```

- (14) `sudo cp /usr/local/hadoop/etc/hadoop/mapred-site.xml.template /usr/local/hadoop/etc/hadoop/mapred-site.xml`

We use super user privileges to copy the file `mapred-site.xml.template` to the new file `mapred-site.xml`.

- (15) `sudo gedit /usr/local/hadoop/etc/hadoop/yarn-site.xml`

We use super user privileges with the text editor `gedit` so as to modify the content of the configuration file `yarn-site.xml`. In our case we fill the (originally empty `<configuration></configuration>`) configuration section with:

```
<configuration>
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
```

```

    </property>
</configuration>

```

- (16) `hadoop`

We launch Hadoop 2.7.1 to ensure it has been installed. In this case we should get the following:

```

Usage: hadoop [--config confdir] [COMMAND | CLASSNAME]
    CLASSNAME                run the class named CLASSNAME
or
    where COMMAND is one of:
    fs                        run a generic filesystem user client
    version                  print the version
    jar <jar>                run a jar file
                             note: please use "yarn jar" to launch
                             YARN applications, not this command.
    checknative [-a|-h]      check native hadoop and compression
    libraries availability
    distcp <srcurl> <desturl> copy file or directories recursively
    archive -archiveName NAME -p <parent path> <src>* <dest>
    create a hadoop archive
    classpath                prints the class path needed to get the
    credential                interact with credential providers
                             Hadoop jar and the required libraries
    daemonlog                get/set the log level for each daemon
    trace                    view and modify Hadoop tracing settings

```

Most commands print help when invoked w/o parameters.

3.4.2 Start a Hadoop Cluster

We start the Hadoop Single Node Cluster with Pseudo-Distributed Operation.

Figure 3.4 presents the script to start the cluster. It can be run from a terminal in our local server.

```

(01) #!/bin/bash
(02) ssh localhost
(03) hdfs namenode -format
(04) start-dfs.sh
(05) start-yarn.sh
(06) hdfs dfs -mkdir /user/
(07) hdfs dfs -mkdir /user/my_HDFS/

```

Figure 3.4: Script Start.Hadoop.Cluster.sh

Next, we present more detailed instructions about the steps followed in the script:

- (01) `#!/bin/bash`

We indicate bin bash as the interpreter being used.

- (02) `ssh localhost`

We connect by ssh to the localhost without a passphrase.

- (03) `hdfs namenode -format`

We ensure HDFS has some format. If the command is successfully executed, then we should get the following, where `***MACHINE_NAME***` represents the name of the server:

```
STARTUP_MSG: Starting NameNode
STARTUP_MSG:  host = ***MACHINE_NAME***/127.0.1.1
STARTUP_MSG:  args = [-format]
STARTUP_MSG:  version = 2.7.1
STARTUP_MSG:  classpath =
STARTUP_MSG:  java = 1.8.0_252
SHUTDOWN_MSG: Shutting down NameNode
                at ***MACHINE_NAME***/127.0.1.1
```

- (04) `start-dfs.sh`

We start the HDFS daemons NameNode and DataNode. If the command is successfully executed, then we should get the following, where `***USER_NAME***` represents the name of the user:

```
Starting namenodes on [localhost]
localhost: starting namenode, logging to /usr/local/hadoop/
logs/hadoop-***USER_NAME***-namenode-***MACHINE_NAME***.out
localhost: starting datanode, logging to /usr/local/hadoop/
logs/hadoop-***USER_NAME***-***MACHINE_NAME***.out
Starting secondary namenodes [0.0.0.0]
0.0.0.0: starting secondarynamenode, logging to /usr/local/
hadoop/logs/hadoop-***USER_NAME***-secondarynamenode-***MACHINE_NAME***
```

- (05) `start-yarn.sh`

We start the YARN Job Scheduler, specifically its daemons ResourceManager and NodeManager. If the command is successfully executed, then we should get the following:

```
starting yarn daemons
starting resourcemanager, logging to /usr/local/hadoop/logs/
yarn-***USER_NAME***-resourcemanager-***MACHINE_NAME***.out
localhost: starting nodemanager, logging to /usr/local/hadoop/
logs/yarn-***USER_NAME***-nodemanager-***MACHINE_NAME***.out
```

- (06) `hdfs dfs -mkdir /user/`

We create the HDFS folder `/user`.

- (06) `hdfs dfs -mkdir /user/my_HDFS/`

We create the HDFS subfolder `/user/my_HDFS/`, where we can place the dataset to be analysed.

3.4.3 Dataset

We use the website [Lipsum \[26\]](#) to generate 4 paragraphs of *Lorem Ipsum* text. We create a dataset `my_dataset` with 4 small text files, each of them containing one of the generated paragraphs. Figure 3.8 presents the content of `my_dataset` in the local file system of our server.

3.4.4 HDFS

Once the cluster is started, we can check the status of the HDFS NameNode at and DataNode at <http://localhost:50070/>

Figures 3.6 and 3.7 show it. As we can see, the folder my_HDFS is empty.

Figure 3.5 presents the script to run a MapReduce or Spark data analysis application on top of HDFS. It can be run from a terminal in our local server.

```
(01) #!/bin/bash
(02) ssh localhost
(03) hdfs dfs -put ./my_dataset/ /user/my_HDFS/my_dataset
(04) # MapReduce or Spark Job Command
(05) hdfs dfs -get /user/my_HDFS/my_result ./
(06) hdfs dfs -rm -r /user/my_HDFS/my_dataset/
(07) hdfs dfs -rm -r /user/my_HDFS/my_result/
```

Figure 3.5: Script data_analysis.sh

Next, we present more detailed instructions about the steps followed in the script:

- (01) #!/bin/bash
We indicate bin bash as the interpreter being used.
- (02) ssh localhost
We connect by ssh to the localhost without a passphrase.
- (03) hdfs dfs -put ./my_dataset/ /user/my_HDFS/my_dataset
We use the HDFS command put to copy my_dataset from the local file system to HDFS. Figure 3.9 shows it.
- (04) # MapReduce or Spark Job Command
We leave out the details of running a MapReduce/Spark Job to the next section. By the moment we assume the MapReduce/Spark Job succeeds, producing the results in the new HDFS folder my_result. Figure 3.10 shows it.
- (05) hdfs dfs -get /user/my_HDFS/my_result ./
We use the HDFS command get to copy my_result from HDFS to our local file system. Figure 3.11 shows it.
- (06) hdfs dfs -rm -r /user/my_HDFS/my_dataset/
- (07) hdfs dfs -rm -r /user/my_HDFS/my_result/
We use the the HDFS command rm to remove the folders my_dataset and my_result from HDFS. By doing so, we restore the content of HDFS to the one presented in Figure 3.7, i.e., ready to run the script again with a new data analysis application.

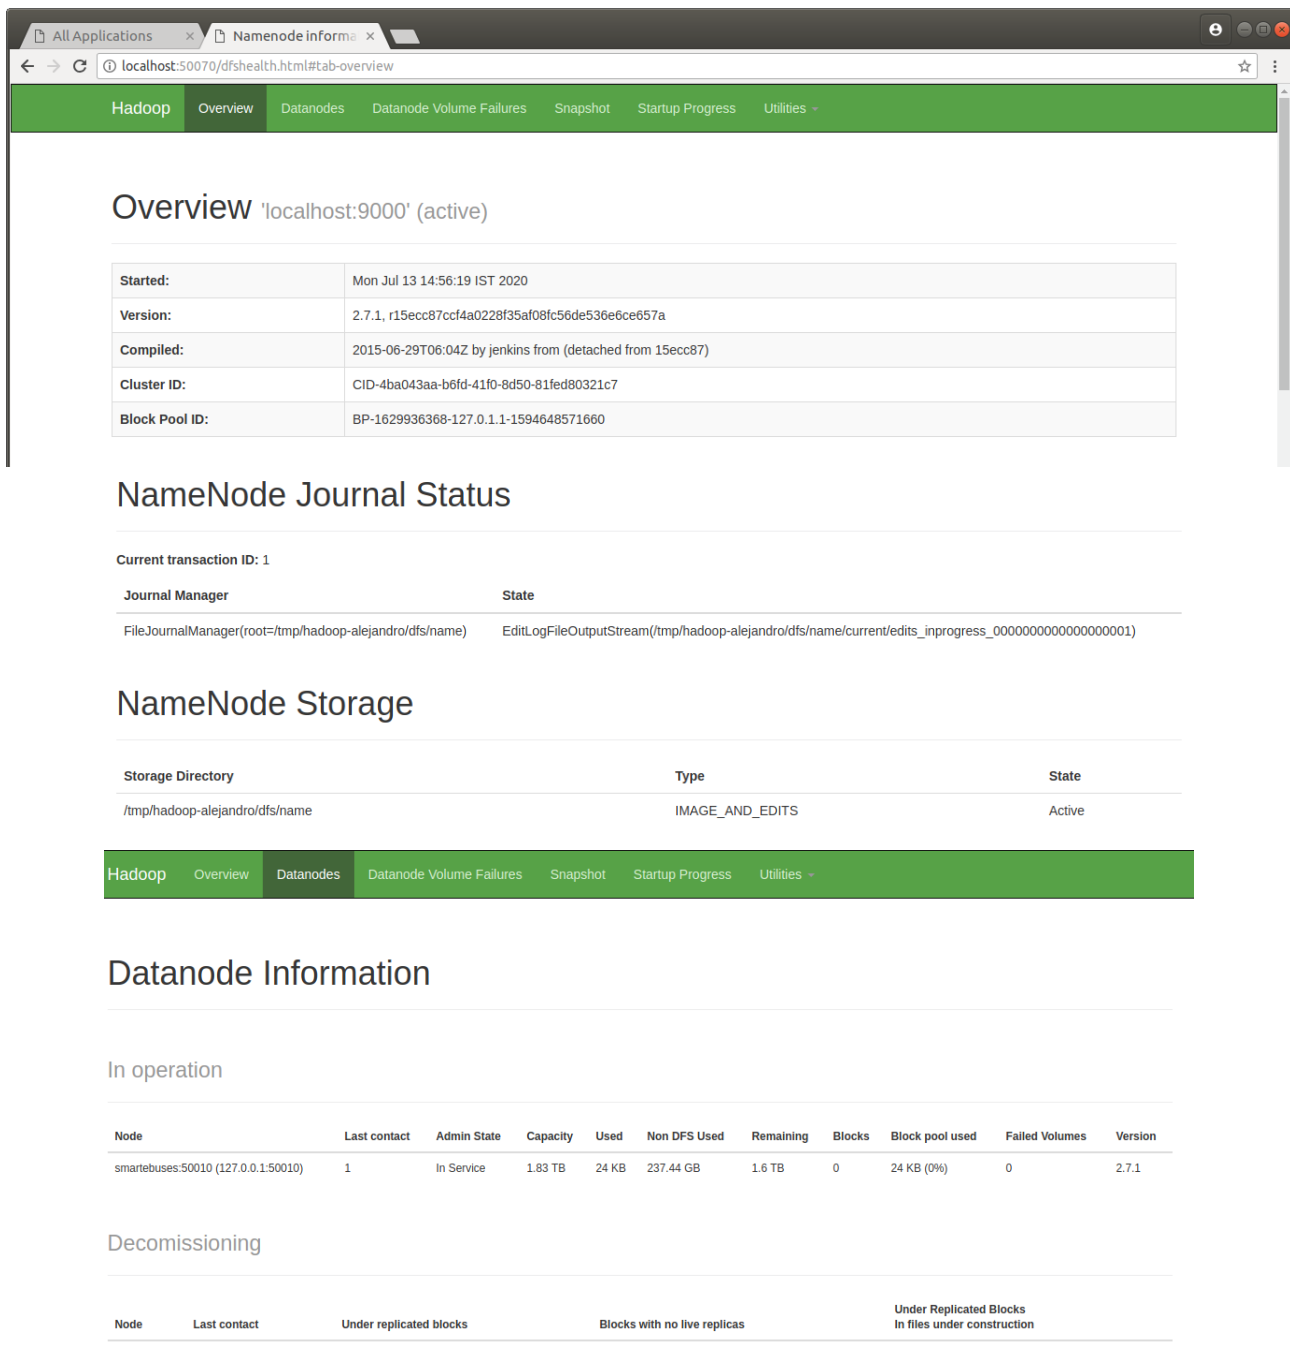


Figure 3.6: Cluster Overview

The figure consists of two screenshots of the Hadoop web interface, specifically the 'Browse Directory' page. Both screenshots have a green navigation bar at the top with links: Hadoop, Overview, Datanodes, Snapshot, Startup Progress, and Utilities.

The top screenshot shows the directory path '/user' in the search bar. Below the search bar is a table with the following data:

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	alejandro	supergroup	0 B	7/13/2020, 3:03:03 PM	0	0 B	my_HDFS

The bottom screenshot shows the directory path '/user/my_HDFS' in the search bar. Below the search bar is an empty table with the same headers as the one above.

Figure 3.7: HDFS

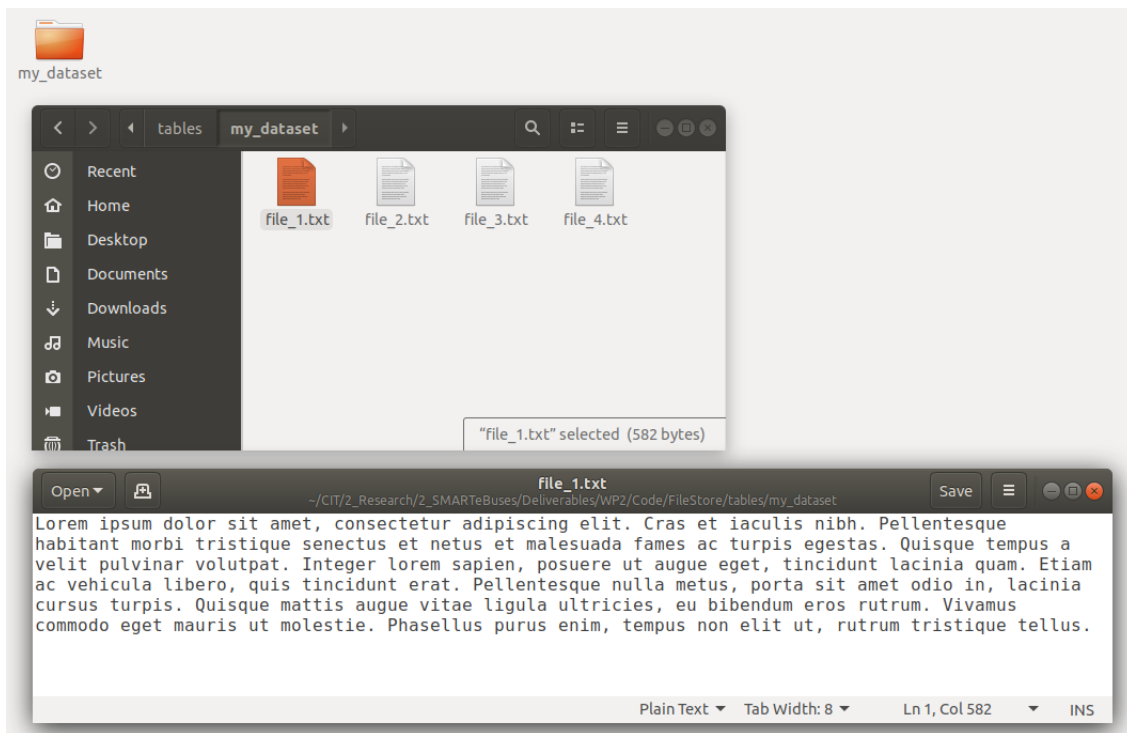


Figure 3.8: Dataset Folder: my_dataset

The top part of the image shows a terminal window with the following commands and output:

```

alejandro@smarte buses: ~/Files
File Edit View Search Terminal Help
alejandro@smarte buses:~/Files$ clear
alejandro@smarte buses:~/Files$ hdfs dfs -put ./my_dataset/ /user/my_HDFS/my_dataset
alejandro@smarte buses:~/Files$

```

The bottom part of the image shows the Hadoop web interface. The breadcrumb path is `/user/my_HDFS/my_dataset`. Below the path is a table listing the files in the directory:

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	alejandro	supergroup	7 B	7/13/2020, 3:08:18 PM	1	128 MB	file_1.txt
-rw-r--r--	alejandro	supergroup	7 B	7/13/2020, 3:08:18 PM	1	128 MB	file_2.txt
-rw-r--r--	alejandro	supergroup	7 B	7/13/2020, 3:08:18 PM	1	128 MB	file_3.txt
-rw-r--r--	alejandro	supergroup	7 B	7/13/2020, 3:08:18 PM	1	128 MB	file_4.txt

A modal window titled "File information - file_1.txt" is open, showing details for Block 0:

- Block ID: 1073741825
- Block Pool ID: BP-1629936368-127.0.1.1-1594648571660
- Generation Stamp: 1001
- Size: 7
- Availability:
 - smarte buses

Figure 3.9: Command put: my_dataset from local file system to HDFS

Hadoop
Overview
Datanodes
Snapshot
Startup Progress
Utilities

Browse Directory

Go!

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	alejandro	supergroup	0 B	13/07/2020, 15:30:11	0	0 B	my_dataset
drwxr-xr-x	alejandro	supergroup	0 B	13/07/2020, 18:19:44	0	0 B	my_result

Hadoop
Overview
Datanodes
Snapshot
Startup Progress
Utilities

Browse Directory

Go!

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	alejandro	supergroup	0 B	13/07/2020, 18:19:44	1	128 MB	_SUCCESS
-rw-r--r--	alejandro	supergroup	1.46 KB	13/07/2020, 18:19:44	1	128 MB	part-00000

Datanodes
Snapshot
Startup Progress
Utilities

File information - part-00000
Download

Block information -- Block 0

Block ID: 1073741849
Block Pool ID: BP-1629936368-127.0.1.1-1594648571660
Generation Stamp: 1025
Size: 14
Availability:

- smartebuses

Close

Figure 3.10: HDFS With New Folder my_result

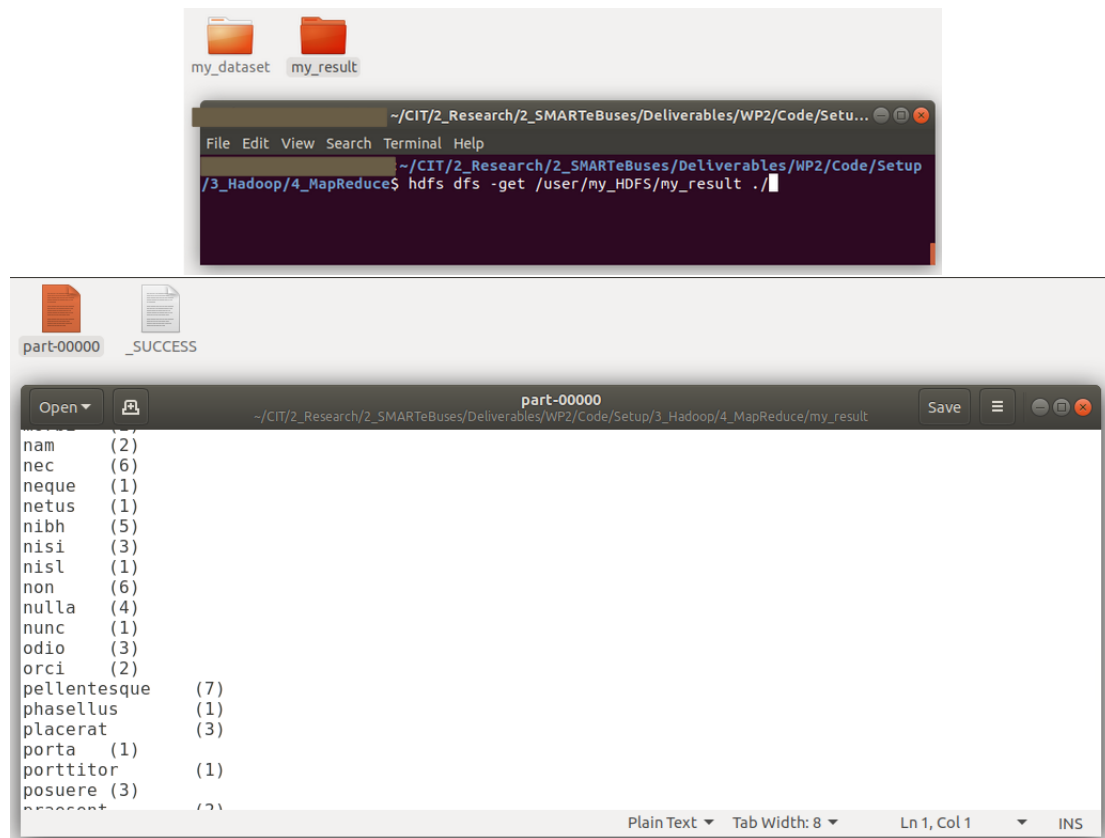


Figure 3.11: Command get: my_result from HDFS to local file system

3.4.5 MapReduce

Once the cluster is started, we can check the status of the Yarn ResourceManager and NodeManager at <http://localhost:8088/>

Figure 3.12 shows it. As we can see, there is no application being run so far.

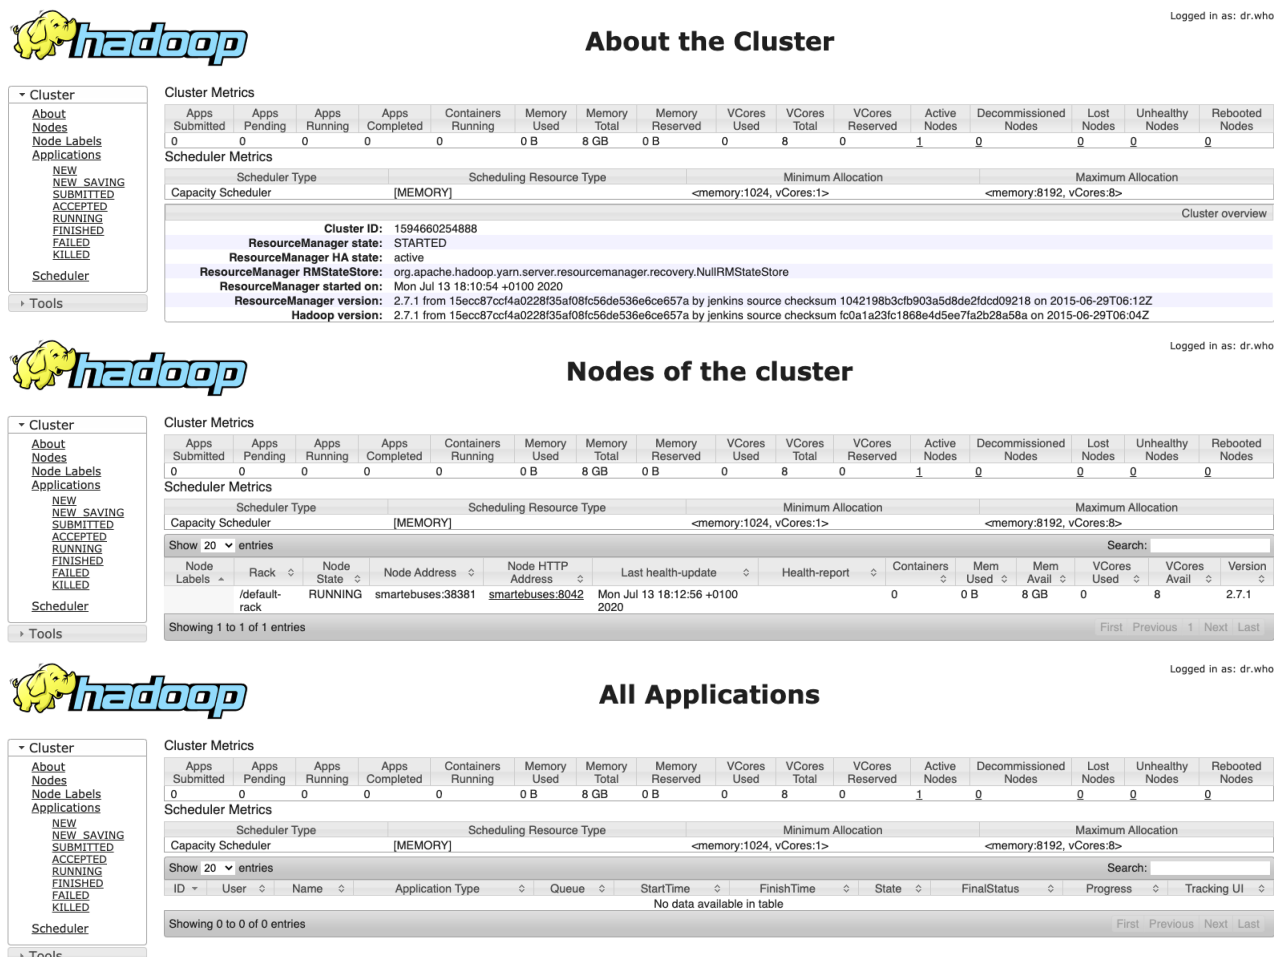



Figure 3.12: ResourceManager Overview



Logged in as: dr.who

All Applications

Cluster

- About
- Nodes
- Node Labels
- Applications
- NEW
- NEW SAVING
- SUBMITTED
- ACCEPTED
- RUNNING
- FINISHED
- FAILED
- KILLED
- Scheduler

Tools

Cluster Metrics															
Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved	Active Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes
1	1	0	0	0	0 B	8 GB	0 B	0	8	0	1	0	0	0	0


Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation
Capacity Scheduler	[MEMORY]	<memory:1024, vCores:1>	<memory:8192, vCores:8>

Show 20 entries

ID	User	Name	Application Type	Queue	StartTime	FinishTime	State	FinalStatus	Progress	Tracking UI
application_1594660254888_0001	alejandra	streamjob4554623490106481759.jar	MAPREDUCE	default	Mon Jul 13 18:19:31 +0100 2020	N/A	ACCEPTED	UNDEFINED		ApplicationMaster

Showing 1 to 1 of 1 entries



Logged in as: dr.who

Application application_1594660254888_0001

Cluster

- About
- Nodes
- Node Labels
- Applications
- NEW
- NEW SAVING
- SUBMITTED
- ACCEPTED
- RUNNING
- FINISHED
- FAILED
- KILLED
- Scheduler

Tools

Kill Application

User: alejandra

Name: streamjob4554623490106481759.jar

Application Type: MAPREDUCE

Application Tags:

YarnApplicationState: ACCEPTED: waiting for AM container to be allocated, launched and register with RM.

FinalStatus Reported by AM: Application has not completed yet.

Started: Mon Jul 13 18:19:31 +0100 2020

Elapsed: 1sec

Tracking URL: ApplicationMaster

Diagnostics:

Application Overview

Application Metrics	
Total Resource Preempted:	<memory:0, vCores:0>
Total Number of Non-AM Containers Preempted:	0
Total Number of AM Containers Preempted:	0
Resource Preempted from Current Attempt:	<memory:0, vCores:0>
Number of Non-AM Containers Preempted from Current Attempt:	0
Aggregate Resource Allocation:	0 MB-seconds, 0 vcore-seconds

Show 20 entries

Attempt ID	Started	Node	Logs
appattemp1_1594660254888_0001_0000001	Mon Jul 13 18:19:31 +0100 2020	http://smarteubuses:8042	Logs

Showing 1 to 1 of 1 entries

Figure 3.13: ResourceManager: MapReduce Application in Progress



Application application_1594660254888_0001

Logged in as: dr.who

Cluster

- About
- Nodes
- Node Labels
- Applications
- NEW
- NEW SAVING
- SUBMITTED
- ACCEPTED
- RUNNING
- FINISHED
- FAILED
- KILLED
- Scheduler

Tools

Kill Application

Application Overview

User:	alejandro
Name:	streamjob4554623490106481759.jar
Application Type:	MAPREDUCE
Application Tags:	
YarnApplicationState:	FINISHED
FinalStatus Reported by AM:	SUCCEEDED
Started:	Mon Jul 13 18:19:31 +0100 2020
Elapsed:	12sec
Tracking URL:	History
Diagnostics:	

Application Metrics

Total Resource Preempted:	<memory:0, vCores:0>
Total Number of Non-AM Containers Preempted:	0
Total Number of AM Containers Preempted:	0
Resource Preempted from Current Attempt:	<memory:0, vCores:0>
Number of Non-AM Containers Preempted from Current Attempt:	0
Aggregate Resource Allocation:	55727 MB-seconds, 33 vcore-seconds

Show 20 entries

Search:

Attempt ID	Started	Node	Logs
appattempt_1594660254888_0001_000001	Mon Jul 13 18:19:31 +0100 2020	http://smarteubuses:8042	Logs

Showing 1 to 1 of 1 entries

First Previous 1 Next Last



All Applications

Logged in as: dr.who

Cluster

- About
- Nodes
- Node Labels
- Applications
- NEW
- NEW SAVING
- SUBMITTED
- ACCEPTED
- RUNNING
- FINISHED
- FAILED
- KILLED
- Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved	Active Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes
1	0	0	1	0	0 B	8 GB	0 B	0	8	0	1	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation
Capacity Scheduler	[MEMORY]	<memory:1024, vCores:1>	<memory:8192, vCores:8>

Show 20 entries

Search:

ID	User	Name	Application Type	Queue	StartTime	FinishTime	State	FinalStatus	Progress	Tracking UI
application_1594660254888_0001	alejandro	streamjob4554623490106481759.jar	MAPREDUCE	default	Mon Jul 13 18:19:31 +0100 2020	Mon Jul 13 18:19:44 +0100 2020	FINISHED	SUCCEEDED		History

Showing 1 to 1 of 1 entries

First Previous 1 Next Last

Figure 3.14: ResourceManager: MapReduce Application Finished

Hadoop
Overview
Datanodes
Snapshot
Startup Progress
Utilities

Browse Directory

Go!

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	alejandro	supergroup	0 B	13/07/2020, 15:30:11	0	0 B	my_dataset
drwxr-xr-x	alejandro	supergroup	0 B	13/07/2020, 18:19:44	0	0 B	my_result

Hadoop
Overview
Datanodes
Snapshot
Startup Progress
Utilities

Browse Directory

Go!

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	alejandro	supergroup	0 B	13/07/2020, 18:19:44	1	128 MB	_SUCCESS
-rw-r--r--	alejandro	supergroup	1.46 KB	13/07/2020, 18:19:44	1	128 MB	part-00000

Figure 3.15: HDFS: MapReduce Result in my_result

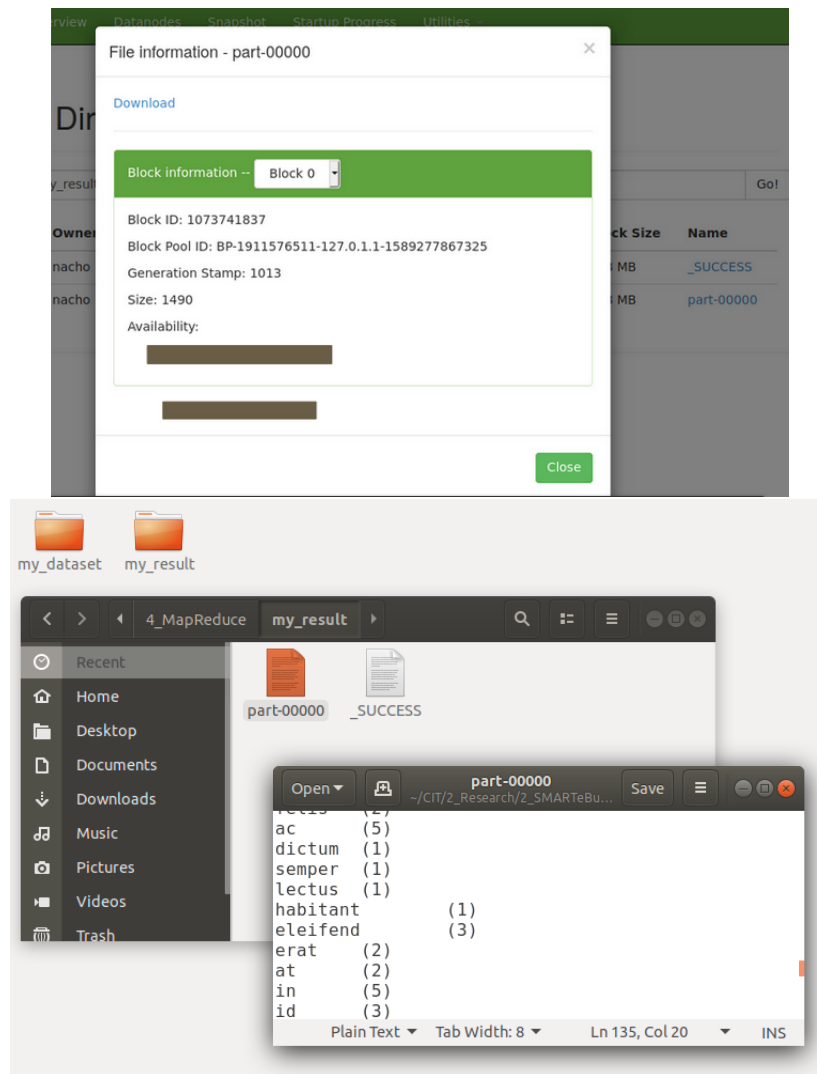


Figure 3.16: HDFS: MapReduce Result Brought Back to Local File System

We edit the command (04) # MapReduce or Spark Job Command from the script `data_analysis.sh` (cf. Figure 3.5) to run our introductory MapReduce application. The MapReduce command is presented below:

```
(04) hadoop jar /usr/local/hadoop/share/hadoop/tools/lib/
      hadoop-streaming-2.7.1.jar \
      -input /user/my_HDFS/my_dataset \
      -output /user/my_HDFS/my_result \
      -mapper ./my_mapper.py \
      -reducer ./my_reducer.py \
      -file ./my_mapper.py \
      -file ./my_reducer.py
```

On it, we use the library Hadoop Streaming [15] to program a MapReduce application where the Map and Reduce stages are specified in the Python files `my_mapper.py` and `my_reducer.py`, resp. Both `my_mapper.py` and `my_reducer.py` are executable scripts, configured to read from the standard input and write to the standard output. The MapReduce application itself is quite simple: it produces as output the new folder `my_result`, containing the word count for the dataset provided in the input folder `my_dataset`.

Figures 3.17, 3.18 and 3.19 present the file `my_mapper.py`.

```
(01) # -----
(02) # IMPORTS
(03) # -----
(04) import sys
(05) import re
```

Figure 3.17: `my_mapper.py`: Import Section

We present more detailed instructions about the steps followed in the script:

- Lines (01)–(05) import the libraries `sys` and `re`. The former is used to redirect the input and output streams to `stdin` and `stdout`, resp. The latter is used to compile regular expressions for processing the content being read.
- Lines (06)–(27) define the function `my_map`. It receives as parameters an input and an output stream. The function reads the whole content provided by the input stream, producing as result a number of text lines to be written to the output stream. In concrete, for each different word being read by the standard input, the function produces a key-value line with the format `word\t(num_appearances)\n`, where `num_appearances` is the amount of times the word `word` has been read.

The function uses a dictionary `my_dict` to collect as keys all words being read, with each word having as its associated key the number of appearances. `my_dict` is initialised in line (11) and it is populated in the for loop of lines (13)–(23). Line (14) ensures we read one line from the standard input at a time. Line (16) splits the text line to the list of words it contains. The loop of lines (18)–(23) processes each word separately. On line (19) we use `re.sub` and `lower` to remove any non-alphabetic character and to lower any upper case letter appearing in the word,

```

(06) # -----
(07) # FUNCTION my_map
(08) # -----
(09) def my_map(my_input_stream,
               my_output_stream
               ):
(10)     # 1. We create a dictionary with all
           the different words in the file
(11)     my_dict = {}
(12)     num_appearances = 1

(13)     # 2. We traverse the file content, to populate my_dict
(14)     for line in my_input_stream:
(15)         # 2.1. We process the line
(16)         word_list = line.split(" ")

(17)         # 2.2. We populate the dictionary
           with the words of the sentence
(18)         for w in word_list:
(19)             my_word = re.sub(r"[^a-zA-Z]", "", w).lower()

(20)             if (my_word in my_dict):
(21)                 my_dict[my_word] =
                     my_dict[my_word] + num_appearances
(22)             else:
(23)                 my_dict[my_word] = num_appearances

(24)     # 3. We write the content of the dict
(25)     for key in my_dict:
(26)         my_str = key + "\t(" + str(my_dict[key]) + ")\n"
(27)         my_output_stream.write(my_str)

```

Figure 3.18: my_mapper.py: my_map Function

resp. Finally, lines (20)–(23) check if the word has already been registered previously. If so, it increments its number of appearances by one. Otherwise, it enters the word in the dictionary with a single appearance.

Once the entire content of the standard input has been read and processed, lines (24)–(27) produce the lines to be written by the output stream. Line (25) traverses the words stored in the dictionary. Line (26) produces the String with the key-value pair associated to the word. Finally, line (27) writes this String to the standard output.

- Lines (31)–(38) define the main entry point for the program. Lines (35)–(36) redirect the input and output streams to `stdin` and `stdout`, resp. Line (38) calls to the aforementioned function `my_map`.

```

(28) # -----
(29) # MAIN
(30) # -----
(31) if __name__ == '__main__':
(32)     # 1. We use as many input arguments as needed
(33)     pass

(34)     # 2. We set the input and output streams
(35)     my_input_stream = sys.stdin
(36)     my_output_stream = sys.stdout

(37)     # 3. We launch the Map program
(38)     my_map(my_input_stream,
              my_output_stream
              )

```

Figure 3.19: my_mapper.py: Main Entry Point

A copy of the file `my_mapper.py` is to be placed on each `DataNode` of the cluster involved in the Map stage, so as to process the subset of `my_dataset` associated to it.

Figures 3.20, 3.21 and 3.22 present the file `my_reducer.py`.

```

(01) # -----
(02) # IMPORTS
(03) # -----
(04) import sys

```

Figure 3.20: my_mapper.py: Import Section

The program is very similar to `my_mapper.py`, so we only highlight the differences:

- Lines (01)–(04) do no longer need to import the library `re`.
- Lines (05)–(26) define the function `my_reduce`. The function has the same responsibility as `my_map`. However, instead of reading from `my_dataset`, the function reads the sorted key-value pairs `word\t(num_appearances)\n` produced by the stage `my_map`. Lines (14)–(17) parse each entry to get its associated word and number of appearances.
- Lines (31)–(38) define the main entry point for the program, calling to the aforementioned function `my_reduce`.

A copy of the file `my_reducer.py` is to be placed on each node of the cluster involved in the Reduce stage, so as to process the subset of the sorted key-value entries associated to it.

Figure 3.13 presents the status of the `ResourceManager` once the command (04) of the script `data_analysis.sh` is launched. As we can see, the MapReduce application is considered to


```
(05) # -----
(06) # FUNCTION my_reduce
(07) # -----
(08) def my_reduce(my_input_stream,
                  my_output_stream
                  ):
(09)     # 1. We create a dictionary with all
        the different words in the file
(10)     my_dict = {}

(11)     # 2. We traverse the file content, to populate my_dict
(12)     for line in my_input_stream:
(13)         # 2.1. We get the info from the line
(14)         line = line.replace("\n", "")
(15)         info = line.split("\t")

(16)         my_word = info[0]
(17)         num_appearances = int(info[1][1:-1])

(18)         # 2.2. We populate the dictionary
            with the words of the sentence
(19)         if (my_word in my_dict):
(20)             my_dict[my_word] =
                my_dict[my_word] + num_appearances
(21)         else:
(22)             my_dict[my_word] = num_appearances

(23)     # 3. We write the content of the dict
(24)     for key in my_dict:
(25)         my_str = key + "\t(" + str(my_dict[key]) + ")\n"
(26)         my_output_stream.write(my_str)
```

Figure 3.21: my_mapper.py: my_map Function

be in progress. Figure 3.14 presents the status once the application finishes. As the execution is successful, the new folder `my_result` is available now in HDFS, with Figure 3.15 showing it. While the content of the files is not directly accessible in HDFS, we can execute the command `get` to bring the folder back to our local file system, so as to explore it (Figure 3.16 shows it). All in all, the MapReduce application finds 135 different words with their associated number of appearances in `my_dataset`.

```

(27) # -----
(28) # MAIN
(29) # -----
(30) if __name__ == '__main__':
(31)     # 1. We use as many input arguments as needed
(32)     pass

(33)     # 2. We set the input and output streams
(34)     my_input_stream = sys.stdin
(35)     my_output_stream = sys.stdout

(36)     # 3. We launch the Map program
(37)     my_map(my_input_stream,
              my_output_stream
              )

```

Figure 3.22: my_mapper.py: Main Entry Point

3.4.6 Stop a Hadoop Cluster

Once we have run our data analysis application, we stop the Hadoop Single Node Cluster with Pseudo-Distributed Operation.

Figure 3.23 presents the script to stop the cluster. It can be run from a terminal in our local server.

```

(01) #!/bin/bash
(02) ssh localhost
(03) hdfs dfs -rm -r /user
(04) stop-yarn.sh
(05) stop-dfs.sh

```

Figure 3.23: Script Stop_Hadoop_Cluster.sh

Next, we present more detailed instructions about the steps followed in the script:

- (01) #!/bin/bash
We indicate bin bash as the interpreter being used.
- (02) ssh localhost
We connect by ssh to the localhost without a passphrase.
- (03) hdfs dfs -rm -r /user
We delete our HDFS folder.
- (04) stop-yarn.sh

We stop the YARN Job Scheduler, specifically its daemons ResourceManager and NodeManager. If the command is successfully executed, then we should get the following:

```
stopping yarn daemons
stopping resourcemanager
localhost: stopping nodemanager
localhost: nodemanager did not stop gracefully after 5 seconds: killing
no proxyserver to stop
```

- (05) stop-dfs.sh

We stop the HDFS daemons NameNode and DataNode. If the command is successfully executed, then we should get the following:

```
Stopping namenodes on [localhost]
localhost: stopping namenode
localhost: stopping datanode
Stopping secondary namenodes [0.0.0.0]
0.0.0.0: stopping secondarynamenode
```

3.5 Spark

In this section we discuss how to install and configure Spark. We also present a Spark Core and a Spark SQL application running on top of the Hadoop cluster described in Section 3.4.

3.5.1 Installing Spark

We install and configure Spark 2.4.5 [27]. While there is a more recent version available (Spark 3.0.0 [28]), this version is still in preview mode and thus it is not stable.

Figure 3.24 presents the script to install and configure Spark 2.4.5. It can be run from a terminal in our local server.

```
(01) #!/bin/bash
(02) wget https://downloads.apache.org/spark/spark-2.4.5/
      spark-2.4.5-bin-hadoop2.7.tgz
(03) tar -xzvf spark-2.4.5-bin-hadoop2.7.tgz
(04) sudo mv spark-2.4.5-bin-hadoop2.7 /usr/local/spark
(05) python3.7 -m pip install pyspark
(06) sudo gedit ~/.bashrc
(07) spark-submit --version
```

Figure 3.24: Script Spark_2_4_5.sh

Next, we present more detailed instructions about the steps followed in the script:

- (01) #!/bin/bash

We indicate bin bash as the interpreter being used.

- (02) `wget https://downloads.apache.org/spark/spark-2.4.5/spark-2.4.5-bin-hadoop2.7.tgz`

We download Spark 2.4.5 as a Gzipped source tarball. Among the different versions, we choose the one that is pre-built for Hadoop 2.7.

- (03) `tar -xzvf spark-2.4.5-bin-hadoop2.7.tgz`

We extract it.

- (04) `sudo mv spark-2.4.5-bin-hadoop2.7 /usr/local/spark`

We move the extracted folder to /usr/local/spark.

- (05) `python3.7 -m pip install pyspark`

We use pip, the package installer for Python, to download `pyspark 2.4.5` [29]. This package is useful as it allows us to develop, debug and run a Spark application locally, using the Python 3.7.7 interpreter and the Python IDE PyCharm [30]. Once the application has been tested, it can be submitted for running in the Hadoop Single Node Cluster with Pseudo-Distributed Operation described in Section 3.4.

- (06) `sudo gedit ~/.bashrc`

We use super user privileges with the text editor `gedit` so as to modify the content of the configuration file `bashrc`. In our case we add:

```
export PATH=$PATH:/usr/local/spark/bin/
export PYSARK_PYTHON=python3.7
```

- (07) `spark-submit --version`

We launch Spark 2.4.5 to ensure it has been installed. In this case we should get the following:

Welcome to

```

/  _/  _  _  _/  /  _
 _\  \/  _  \/  _  \/  _/  '  _/
/_  _/  .  _/\  _  _/  _/  _\  _
      /  _/

```

version 2.4.5

Using Scala version 2.11.12, OpenJDK 64-Bit Server VM, 1.8.0_252
Branch HEAD

Compiled by user centos on 2020-02-02T19:38:06Z

Revision cee4ecbb16917fa85f02c635925e2687400aa56b

```
Url https://gitbox.apache.org/repos/asf/spark.git
```

Type --help for more information.

3.5.2 Spark Core

In Section 3.4 we started the cluster and run a MapReduce application on it. Figure 3.14 showed the status of the Yarn ResourceManager and NodeManager (available at <http://localhost:8088/>) including such MapReduce application.

We edit now the command (04) `# MapReduce or Spark Job Command` from the script `data_analysis.sh` (cf. Figure 3.5) to run our introductory Spark Core application. The Spark Core command is presented below:

```
(04) spark-submit \
```

```
--master yarn --deploy-mode cluster \
./my_Spark_Core_example.py \
/user/my_HDFS/my_dataset \
/user/my_HDFS/my_result
```

On it, we use `spark-submit` to launch the Spark Core application in the Hadoop Single Node Cluster with Pseudo-Distributed Operation. The line `--master yarn --deploy-mode cluster` specifies Yarn to be the ResourceManager handling the application. The line `./my_Spark_Core_example.py` specifies the Python file containing the Spark Core application. Finally, the lines `/user/my_HDFS/my_dataset` and `/user/my_HDFS/my_result` are the first and second parameter of the Python program, resp.

The functionality of the Spark Core application is equivalent to the one of the MapReduce application presented in Section 3.4: it produces as output the new folder `my_result`, containing the word count for the dataset provided in the input folder `my_dataset`.

Figures 3.25, 3.26 and 3.27 present the file `my_Spark_Core_example.py`.

```
(01) # -----
(02) # IMPORTS
(03) # -----
(04) import pyspark
(05) import sys
(06) import re
```

Figure 3.25: `my_Spark_Core_example.py`: Import Section

We present more detailed instructions about the steps followed in the script:

- Lines (01)–(06) import the libraries `pyspark`, `sys` and `re`. The library `pyspark` allows us to use the Spark Core API. The library `sys` is used to pass the input and output directories to the program. The library `re` is used to compile regular expressions for processing the content being read.
- Lines (07)–(20) define the function `my_spark_core_job`. It receives as parameters the `SparkContext` (`sc`) and the aforementioned input and output directories (`my_dataset_dir` and `my_result_dir`, resp). The function processes the dataset in `my_dataset_dir` to produce the new folder `my_result_dir` with its word count.

Line (12) applies the creator operation `textFile` (with the folder `my_dataset_dir` as parameter) to load the dataset into `inputRDD`, an RDD of String items. In particular, each line of text of the dataset file turns into one item of such RDD. Figure 3.8 showed the content of `file_1.txt` of `my_dataset`. Figure 3.28 presents how the first 2 lines of the file translate into 2 items of `inputRDD`.

Line (14) applies the transformation operation `flatMap` (with the lambda function `lambda line: line.split(" ")` as parameter) on `inputRDD` to produce the new RDD `all_wordsRDD` of String items. In particular, each item (text line) of `inputRDD` is to be split into a list with the words it contains, and then this word list is exploded, producing

```

(07) # -----
(08) # FUNCTION my_spark_core_job
(09) # -----
(10) def my_spark_core_job(sc, my_dataset_dir, my_result_dir):
(11)     # 1. Operation C1: textFile
(12)     inputRDD = sc.textFile(my_dataset_dir)

(13)     # 2. Operation T1: flatMap
(14)     all_wordsRDD =
            inputRDD.flatMap(lambda line: line.split(" "))

(15)     # 3. Operation T2: map
(16)     clean_wordsRDD =
            all_wordsRDD.map(lambda w:
                (re.sub(r"[^a-zA-Z]", "", w).lower(),
                 1
                )
            )

(17)     # 4. Operation T3: reduceByKey
(18)     solutionRDD =
            clean_wordsRDD.reduceByKey(lambda x, y: x + y)

(19)     # 5. Operation A1: saveAsTextFile
(20)     solutionRDD.saveAsTextFile(my_result_dir)

```

Figure 3.26: my_Spark_Core_example.py: my_spark_core_job Function

one item per word. Figure 3.29 presents the items produced in `all_wordsRDD` for the first item of `inputRDD` showed in Figure 3.28.

Line (16) applies the transformation operation `map` (with the lambda function `lambda w: (re.sub(r"[^a-zA-Z]", "", w).lower(), 1)` as parameter) on `all_wordsRDD` to produce the new RDD `clean_wordsRDD` of tuple (String, int) items. In particular, each item (word) of `all_wordsRDD` is to see removed any non-alphanumeric character; besides that, any upper-case letter is turned into its equivalent lower-case letter. The function returns the tuple (word, 1), registering one appearance of the word. Figure 3.30 presents the items produced in `clean_wordsRDD` for the subset of items of `all_wordsRDD` showed in Figure 3.29.

Line (18) applies the transformation operation `reduceByKey` (with the lambda function `lambda x, y: x + y` as parameter) on `clean_wordsRDD` to produce the new RDD `solutionRDD` of tuple (String, int) items. In particular, all items containing the same key (word) are aggregated (adding the number of appearances). Figure 3.31 presents the content of `solutionRDD` for the words showed in Figure 3.30.

Line (20) applies the action operation `saveAsTextFile` (with the folder `my_result_dir`

```

(21) # -----
(22) # MAIN
(23) # -----
(24)     # 1. We use as many input arguments as needed
(25)     my_dataset_dir = "/FileStore/tables/my_dataset/"
(26)     my_result_dir = "/FileStore/tables/my_result/"

(27)     if (len(sys.argv) > 1):
(28)         my_dataset_dir = sys.argv[1]
(29)         my_result_dir = sys.argv[2]

(30)     # 2. We configure the Spark Context
(31)     sc = pyspark.SparkContext.getOrCreate()
(32)     sc.setLogLevel('WARN')
(33)     print("\n\n\n")

(34)     # 3. We call to our main function
(35)     my_spark_core_job(sc, my_dataset_dir, my_result_dir)

```

Figure 3.27: my_Spark_Core_example.py: Main Entry Point

as parameter) to store the content of `solutionRDD` into it. The result folder `my_result_dir` contains as many files as partitions are in `solutionRDD`. In particular, each item of the RDD is stored as a text line. Figures 3.34 and 3.35 present the files and content produced by `solutionRDD` when the program is executed in the Hadoop cluster.

- Lines (21)–(35) define the main entry point for the program. Lines (25)–(26) specify the input and output directories. When testing the program locally, we run the program in PyCharm and assign the dataset folder `my_dataset` of the local file system. When running the program in the Hadoop cluster we use lines (27)–(29) to set the input and output directories to the HDFS folders `/user/my_HDFS/my_dataset` and `/user/my_HDFS/my_result` (as described in the command (04) `spark submit` of the script `data_analysis.sh` for launching the program). Lines (31)–(33) create the `SparkContext`, configuring how verbose it should be on reporting the status when running the application. Finally, line (35) calls to the aforementioned function `my_spark_core_job`.

A copy of the file `my_Spark_Core_example.py` is to be placed on the `DataNode` the cluster running the Spark Driver process.

Figure 3.32 presents the status of the `ResourceManager` once the command (04) of the script `data_analysis.sh` is launched. As we can see, the Spark Core application is considered to be in progress. Figure 3.33 presents the status once the application finishes. As the execution is successful, the new folder `my_result` is available now in HDFS, with Figure 3.34 showing it. While the content of the files is not directly accessible in HDFS, we can execute the command `get` to bring the folder back to our local file system, so as to explore it (Figure 3.35 shows it). All in all, the Spark Core application finds 135 different words with their associated number of appearances in `my_dataset`.

```

Lorem ipsum dolor sit amet, elit. Cras et nibh. Pellentesque\n
habitant morbi tristique senectus et netus et malesuada fames ac
turpis egestas. Quisque tempus a\n
...

```

Figure 3.28: inputRDD Content

```

Lorem
ipsum
dolor
sit
amet,
elit.
Cras
et
nibh.
Pellentesque\n
...

```

Figure 3.29: all_wordsRDD Content

```

('lorem', 1)
('ipsum', 1)
('dolor', 1)
('sit', 1)
('amet', 1)
('elit', 1)
('cras', 1)
('et', 1)
('nibh', 1)
('pellentesque', 1)
...

```


Figure 3.30: clean_wordsRDD Content

```

('lorem', 4)
('ipsum', 2)
('dolor', 4)
('sit', 6)
('amet', 6)
('elit', 2)
('cras', 3)
('et', 4)
('nibh', 5)
('pellentesque', 7)
...

```

Figure 3.31: RDDs Content



Logged in as: dr.who

All Applications

Cluster

- About
- Nodes
- Node Labels
- Applications
- NEW
- NEW SAVING
- SUBMITTED
- ACCEPTED
- RUNNING
- FINISHED
- FAILED
- KILLED
- Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved	Active Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes
1	0	1	0	1	2 GB	8 GB	0 B	1	8	0	1	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation
Capacity Scheduler	[MEMORY]	<memory:1024, vCores:1>	<memory:8192, vCores:8>

Show 20 entries

ID	User	Name	Application Type	Queue	StartTime	FinishTime	State	FinalStatus	Progress	Tracking UI
application_1594657479419_0001	alejandro	my_Spark_Core_example.py	SPARK	default	Mon Jul 13 17:29:54 +0100 2020	N/A	RUNNING	UNDEFINED		ApplicationMaster

Showing 1 to 1 of 1 entries



Logged in as: dr.who

Application application_1594657479419_0001

Cluster

- About
- Nodes
- Node Labels
- Applications
- NEW
- NEW SAVING
- SUBMITTED
- ACCEPTED
- RUNNING
- FINISHED
- FAILED
- KILLED
- Scheduler

Tools

Kill Application

User: alejandro

Name: my_Spark_Core_example.py

Application Type: SPARK

Application Tags:

YarnApplicationState: RUNNING: AM has registered with RM and started running.

FinalStatus Reported by AM: Application has not completed yet.

Started: Mon Jul 13 17:29:54 +0100 2020

Elapsed: 7sec

Tracking URL: ApplicationMaster

Diagnostics:

Application Metrics

Total Resource Preempted: <memory:0, vCores:0>

Total Number of Non-AM Containers Preempted: 0

Total Number of AM Containers Preempted: 0

Resource Preempted from Current Attempt: <memory:0, vCores:0>

Number of Non-AM Containers Preempted from Current Attempt: 0


Aggregate Resource Allocation: 12874 MB-seconds, 6 vcore-seconds

Show 20 entries

Attempt ID	Started	Node	Logs
appattempt_1594657479419_0001_000001	Mon Jul 13 17:29:54 +0100 2020	http://smarte buses:8042	Logs

Showing 1 to 1 of 1 entries

Figure 3.32: ResourceManager: Spark Core Application in Progress



Application application_1594657479419_0001

Logged in as: dr.who

Cluster

- About
- Nodes
- Node Labels
- Applications
- NEW
- NEW SAVING
- SUBMITTED
- ACCEPTED
- RUNNING
- FINISHED
- FAILED
- KILLED
- Scheduler

Tools

Kill Application

Application Overview

User: alejandro
Name: my_Spark_Core_example.py
Application Type: SPARK
Application Tags:
YarnApplicationState: FINISHED
FinalStatus Reported by AM: SUCCEEDED
Started: Mon Jul 13 17:29:54 +0100 2020
Elapsed: 14sec
Tracking URL: History
Diagnostics:

Application Metrics

Total Resource Preempted: <memory:0, vCores:0>
Total Number of Non-AM Containers Preempted: 0
Total Number of AM Containers Preempted: 0
Resource Preempted from Current Attempt: <memory:0, vCores:0>
Number of Non-AM Containers Preempted from Current Attempt: 0
Aggregate Resource Allocation: 76312 MB-seconds, 37 vcore-seconds

Show 20 entries

Attempt ID

Started


Node

Logs

appattempt_1594657479419_0001_000001
Mon Jul 13 17:29:54 +0100 2020
http://smarteubuses:8042
Logs

Showing 1 to 1 of 1 entries

First Previous 1 Next Last



All Applications

Logged in as: dr.who

Cluster

- About
- Nodes
- Node Labels
- Applications
- NEW
- NEW SAVING
- SUBMITTED
- ACCEPTED
- RUNNING
- FINISHED
- FAILED
- KILLED
- Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved	Active Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes
1	0	0	1	0	0 B	8 GB	0 B	0	8	0	1	0	0	0	0

Scheduler Metrics

Scheduler Type

Scheduling Resource Type

Minimum Allocation

Maximum Allocation

Capacity Scheduler
(MEMORY)
<memory:1024, vCores:1>
<memory:8192, vCores:8>

Show 20 entries

ID

User

Name

Application Type

Queue

StartTime

FinishTime

State

FinalStatus

Progress

Tracking UI

application_1594657479419_0001
alejandro
my_Spark_Core_example.py
SPARK
default
Mon Jul 13 17:29:54 +0100 2020
Mon Jul 13 17:30:09 +0100 2020
FINISHED
SUCCEEDED

History

Showing 1 to 1 of 1 entries

First Previous 1 Next Last

Figure 3.33: ResourceManager: Spark Core Application Finished

Hadoop Overview Datanodes Snapshot Startup Progress Utilities -							
Browse Directory							
/user/my_HDFS/my_result							Go!
Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	alejandro	supergroup	0 B	13/07/2020, 17:49:09	1	128 MB	_SUCCESS
-rw-r--r--	alejandro	supergroup	606 B	13/07/2020, 17:49:09	1	128 MB	part-00000
-rw-r--r--	alejandro	supergroup	427 B	13/07/2020, 17:49:09	1	128 MB	part-00001
-rw-r--r--	alejandro	supergroup	362 B	13/07/2020, 17:49:09	1	128 MB	part-00002
-rw-r--r--	alejandro	supergroup	500 B	13/07/2020, 17:49:09	1	128 MB	part-00003

Figure 3.34: HDFS: Spark Core Result in my_result

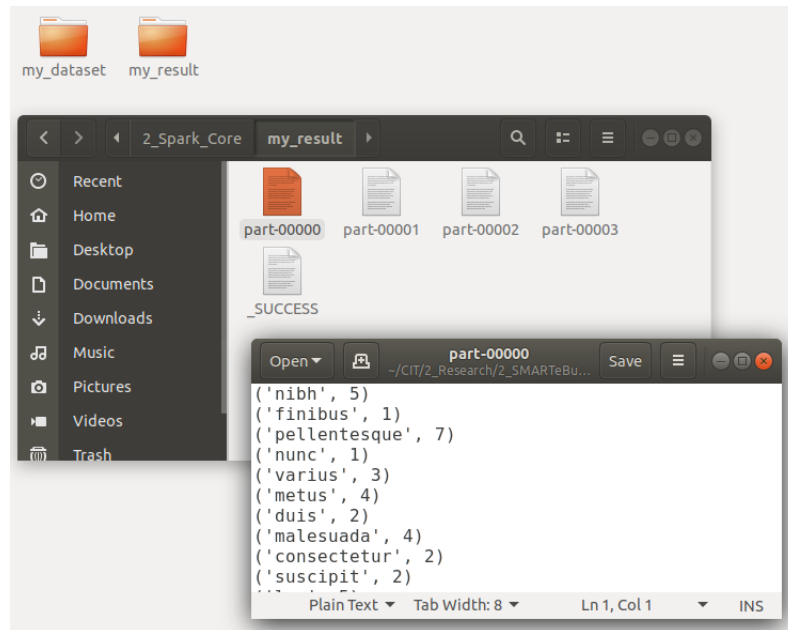


Figure 3.35: HDFS: Spark Core Result Brought Back to Local File System

3.5.3 Spark SQL

We edit now the command (04) # MapReduce or Spark Job Command from the script `data_analysis.sh` (cf. Figure 3.5) to run our introductory Spark SQL application. The Spark SQL command is presented below:

```
(04) spark-submit \
    --master yarn --deploy-mode cluster \
    ./my_Spark_SQL_example.py \
    /user/my_HDFS/my_dataset \
    /user/my_HDFS/my_result
```

As we can see, the command is the same as for the Spark Core application, it only changes the name of the Python program. And so it is the functionality of the Spark SQL application equivalent to the ones of the MapReduce and the Spark Core applications: it produces as output the new folder `my_result`, containing the word count for the dataset provided in the input folder `my_dataset`.

Figures 3.36, 3.37, 3.38 and 3.39 present the file `my_Spark_SQL_example.py`.

```
(01) # -----
(02) # IMPORTS
(03) # -----
(04) import pyspark
(05) import pyspark.sql.functions
(06) import pyspark.sql.types
(07) import sys
```

Figure 3.36: `my_Spark_SQL_example.py`: Import Section

The program is very similar to `my_spark_core_example.py`, so we only highlight the differences:

- Lines (01)–(07) import the additional libraries `pyspark.sql.functions` and `pyspark.sql.types` to use the Spark SQL API.
- Lines (08)–(36) define the function `my_spark_sql_job`. It receives as parameters the `SparkSession` (`spark`) and the input and output directories (`my_dataset_dir` and `my_result_dir`, resp). The function processes the dataset in `my_dataset_dir` to produce the new folder `my_result_dir` with its word count.

Line (12) defines the schema `my_schema` for the dataset, matching each line of text of the dataset to a Row object with a single column `line`, of type `String`.

Line (14) applies the creator operation `read` (with the folder `my_dataset_dir` and `my_schema` as parameters) to load the dataset into `inputDF`, a `DataFrame` of Row objects with one column `line` of type `String`. Figure 3.8 showed the content of `file_1.txt` of `my_dataset`. Figure 3.40 presents how the first 2 lines of the file translate into 2 Row items of `inputDF`.

Lines (16)–(18) apply the transformation operations `withColumn` and `drop` on `inputDF` to produce the new DF `sentenceDF` of Row objects with one column `words_list` of type

list of String. In particular, in line (17) the operation `withColumn` creates the new column `words_list` by splitting the String of `line` into its words. Then, in line (18) the operation `drop` removes the column `line`, for it to not appear in the generated `sentenceDF`. Figure 3.41 presents the Row produced in `sentenceDF` for the first item of `inputDF` showed in Figure 3.40.

Lines (20)–(30) repeat the application of `withColumn` and `drop` over 3 consecutive DataFrames, to shape the content of its unique column to the desired format by: exploding the list of words into a single Row per word (line 21), removing any non-alphanumeric character (line 25) and turning any upper-case character into its equivalent lower-case one (line 29). Likewise, lines (22), (26) and (30) remove any intermediate column being produced. Figures 3.42, 3.43 and 3.44 show the content of these 3 DataFrames, which end up producing the new DataFrame `lowerDF`.

Lines (32)–(33) apply the transformation operation `groupBy` on `lowerDF` to produce the new DF `solutionDF` of Row objects with two column: `word` (of type String) and `count(word)` (of type Integer). In particular, in line (33) the operation `groupBy` specifies the aggregation of Column `word`, counting the appearances for each group. Figure 3.45 shows the content of `solutionDF` for the Rows showed in Figure 3.44.

Lines (35)–(36) apply the action operation `write` (with the folder `my_result_dir` as parameter) to store the content of `solutionDF` into it. Figures 3.48 and 3.49 present the files and content produced by `solutionDF` when the program is executed in the Hadoop cluster.

- Lines (37)–(51) define the main entry point for the program. In particular, line (47) creates the `SparkSession`. Finally, line (51) calls to the aforementioned function `my_spark_sql_job`.

A copy of the file `my_Spark_SQL_example.py` is to be placed on the DataNode the cluster running the Spark Driver process.

Figure 3.46 presents the status of the `ResourceManager` once the command (04) of the script `data_analysis.sh` is launched. As we can see, the Spark SQL application is considered to be in progress. Figure 3.47 presents the status once the application finishes. As the execution is successful, the new folder `my_result` is available now in HDFS, with Figure 3.48 showing it. While the content of the files is not directly accessible in HDFS, we can execute the command `get` to bring the folder back to our local file system, so as to explore it (Figure 3.49 shows it). All in all, the Spark SQL application finds 135 different words with their associated number of appearances in `my_dataset`.

```

(08) # -----
(09) # FUNCTION my_spark_sql_job
(10) # -----
(11) def my_spark_sql_job(spark, my_dataset_dir, my_result_dir):
(11) # 1. We define the Schema of our DF.
(12)     my_schema = pyspark.sql.types.StructType(
                        [pyspark.sql.types.StructField("line",
                        pyspark.sql.types.StringType(), True)
                        ]
                        )

(13)     # 2. Operation C1: Load DataFrame
(14)     inputDF = spark.read.format("csv") \
                        .option("delimiter", ";") \
                        .option("quote", "") \
                        .option("header", "false") \
                        .schema(my_schema) \
                        .load(my_dataset_dir)

(15)     # 3. Operation T1: split
(16)     sentenceDF = inputDF \
(17)         .withColumn("words_list",
                        pyspark.sql.functions.split(
                        pyspark.sql.functions.col("line"),
                        " "
                        )
                        ) \
(18)         .drop("line")

(19)     # 4. Operation T2: explode
(20)     wordsDF = sentenceDF \
(21)         .withColumn("draft_word",
                        pyspark.sql.functions.explode(
                        pyspark.sql.functions.col("words_list")
                        )
                        ) \
(22)         .drop("words_list")

```

Figure 3.37: my_Spark_SQL_example.py: my_spark_sql_job Function

```
(23)      # 5. Operation T3: regexp_replace
(24)      cleanDF = wordsDF \
(25)          .withColumn("clean_word",
                        pyspark.sql.functions.regexp_replace(
                            "draft_word",
                            r"^[a-zA-Z]",
                            ""
                        ) \
(26)          .drop("draft_word")

(27)      # 6. Operation T4: lower
(28)      lowerDF = cleanDF \
(29)          .withColumn("word",
                        pyspark.sql.functions.lower(
                            pyspark.sql.functions.col("clean_word")
                        ) \
(30)          .drop("clean_word")

(31)      # 7. Operation T5: GroupBy
(32)      solutionDF = lowerDF \
(33)          .groupBy(["word"]).agg({"word": "count"})

(34)      # 8. Operation A1: We save the results
(35)      solutionDF.write.format("csv") \
(36)          .save(my_result_dir)
```

Figure 3.38: my_Spark_SQL_example.py: my_spark_sql_job Function


```
(37) # -----
(38) # MAIN
(39) # -----
(40)     # 1. We use as many input arguments as needed
(41)     my_dataset_dir = "/FileStore/tables/my_dataset/"
(42)     my_result_dir = "/FileStore/tables/my_result/"

(43)     if (len(sys.argv) > 1):
(44)         my_dataset_dir = sys.argv[1]
(45)         my_result_dir = sys.argv[2]

(46)     # 2. We configure the Spark Session
(47)     spark = pyspark.sql.Session.builder.getOrCreate()
(48)     spark.sparkContext.setLogLevel('WARN')
(49)     print("\n\n\n")

(50)     # 3. We call to our main function
(51)     my_spark_sql_job(spark, my_dataset_dir, my_result_dir)
```

Figure 3.39: my_Spark_SQL_example.py: Main Entry Point

```
+-----+
|          line|
+-----+
|Lorem ipsum dolor...|
|habitant morbi tr...|
|...                |
+-----+
```

Figure 3.40: inputDF Content

```
+-----+
|      words_list|
+-----+
|[Lorem, ipsum., d...|
|...                |
+-----+
```

Figure 3.41: sentenceDF Content

```
+-----+
|      draft_word|
+-----+
|Lorem           |
|ipsum.          |
|dolor           |
|...             |
+-----+
```

Figure 3.42: wordsDF Content

```

+-----+
|          clean_word|
+-----+
|Lorem                |
|ipsum                |
|dolor                |
|...                  |
+-----+

```

Figure 3.43: cleanDF Content

```

+-----+
|          word|
+-----+
|lorem         |
|ipsum         |
|dolor         |
|...           |
+-----+

```


Figure 3.44: lowerDF Content

```

+-----+-----+
|          word|count (word) |
+-----+-----+
|lorem         |          4 |
|ipsum         |          2 |
|dolor         |          4 |
|...           |          ...|
+-----+-----+

```

Figure 3.45: solutionDF Content



Logged in as: dr.who

All Applications

Cluster

- About
- Nodes
- Node Labels
- Applications
- NEW
- NEW SAVING
- SUBMITTED
- ACCEPTED
- RUNNING
- FINISHED
- FAILED
- KILLED
- Scheduler

Tools

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved	Active Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes
3	0	1	2	1	2 GB	8 GB	0 B	1	8	0	1	0	0	0	0


Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation
Capacity Scheduler	[MEMORY]	<memory:1024, vCores:1>	<memory:8192, vCores:8>

Show 20 entries

ID	User	Name	Application Type	Queue	StartTime	FinishTime	State	FinalStatus	Progress	Tracking UI
application_1594657479419_0003	alejandro	my_Spark_SQL_example.py	SPARK	default	Mon Jul 13 17:55:57 +0100 2020	N/A	ACCEPTED	UNDEFINED		ApplicationMaster
application_1594657479419_0002	alejandro	my_Spark_Core_example.py	SPARK	default	Mon Jul 13 17:48:56 +0100 2020	Mon Jul 13 17:49:10 +0100 2020	FINISHED	SUCCEEDED		History
application_1594657479419_0001	alejandro	my_Spark_Core_example.py	SPARK	default	Mon Jul 13 17:29:54 +0100 2020	Mon Jul 13 17:30:09 +0100 2020	FINISHED	SUCCEEDED		History

Showing 1 to 3 of 3 entries



Logged in as: dr.who

Application application_1594657479419_0003

Cluster

- About
- Nodes
- Node Labels
- Applications
- NEW
- NEW SAVING
- SUBMITTED
- ACCEPTED
- RUNNING
- FINISHED
- FAILED
- KILLED
- Scheduler

Tools

Kill Application

User: alejandro

Name: my_Spark_SQL_example.py

Application Type: SPARK

Application Tags:

YarnApplicationState: RUNNING: AM has registered with RM and started running.

FinalStatus Reported by AM: Application has not completed yet.

Started: Mon Jul 13 17:55:57 +0100 2020

Elapsed: 4sec

Tracking URL: ApplicationMaster

Diagnostics:

Application Overview

Application Metrics	
Total Resource Preempted:	<memory:0, vCores:0>
Total Number of Non-AM Containers Preempted:	0
Total Number of AM Containers Preempted:	0
Resource Preempted from Current Attempt:	<memory:0, vCores:0>
Number of Non-AM Containers Preempted from Current Attempt:	0
Aggregate Resource Allocation:	5373 MB-seconds, 2 vcore-seconds

Show 20 entries

Attempt ID	Started	Node	Logs
appattemp1_1594657479419_0003_000001	Mon Jul 13 17:55:57 +0100 2020	http://smarteubuses.8042	Logs

Showing 1 to 1 of 1 entries

Figure 3.46: ResourceManager: Spark SQL Application in Progress



All Applications

Logged in as: dr.who

Cluster

About Nodes

Node Labels

Applications

NEW

NEW SAVING

SUBMITTED

ACCEPTED

RUNNING

FINISHED

FAILED

KILLED

Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved	Active Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes
3	0	0	3	0	0 B	8 GB	0 B	0	8	0	1	0	0	0	0

Scheduler Metrics

Scheduler Type		Scheduling Resource Type		Minimum Allocation				Maximum Allocation					
Capacity Scheduler		[MEMORY]		<memory:1024, vCores:1>				<memory:8192, vCores:8>					
Show 20 entries												Search:	
ID	User	Name	Application Type	Queue	StartTime	FinishTime	State	FinalStatus	Progress	Tracking UI			
application_1594657479419_0003	alejandro	my_Spark_SQL_example.py	SPARK	default	Mon Jul 13 17:55:57 +0100 2020	Mon Jul 13 17:56:19 +0100 2020	FINISHED	SUCCEEDED		History			
application_1594657479419_0002	alejandro	my_Spark_Core_example.py	SPARK	default	Mon Jul 13 17:48:56 +0100 2020	Mon Jul 13 17:49:10 +0100 2020	FINISHED	SUCCEEDED		History			
application_1594657479419_0001	alejandro	my_Spark_Core_example.py	SPARK	default	Mon Jul 13 17:29:54 +0100 2020	Mon Jul 13 17:30:09 +0100 2020	FINISHED	SUCCEEDED		History			
Showing 1 to 3 of 3 entries												First Previous 1 Next Last	

Showing 1 to 3 of 3 entries

First Previous 1 Next Last



Application application_1594657479419_0002

Logged in as: dr.who

Cluster

About

Nodes

Node Labels

Applications

NEW

NEW SAVING

SUBMITTED

ACCEPTED

RUNNING

FINISHED

FAILED

KILLED

Scheduler

Tools

Kill Application

User: alejandro

Name: my_Spark_Core_example.py

Application Type: SPARK

Application Tags:

YarnApplicationState: FINISHED

FinalStatus Reported by AM: SUCCEEDED

Started: Mon Jul 13 17:48:56 +0100 2020

Elapsed: 13sec

Tracking URL: History

Diagnostics:

Application Overview

Total Resource Preempted: <memory:0, vCores:0>

Total Number of Non-AM Containers Preempted: 0

Total Number of AM Containers Preempted: 0

Resource Preempted from Current Attempt: <memory:0, vCores:0>

Number of Non-AM Containers Preempted from Current Attempt: 0

Aggregate Resource Allocation: 75915 MB-seconds, 37 vcore-seconds

Application Metrics

Show 20 entries

Attempt ID

Started

Node

Logs

appattempt_1594657479419_0002_000001

Mon Jul 13 17:48:56 +0100 2020

http://smarte buses:8042

Logs

Showing 1 to 1 of 1 entries

First Previous 1 Next Last

Showing 1 to 1 of 1 entries

First Previous 1 Next Last

Figure 3.47: ResourceManager: Spark SQL Application Finished

Hadoop Overview Datanodes Snapshot Startup Progress Utilities -							
Browse Directory							
/user/my_HDFS/my_result							Go!
Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	alejandro	supergroup	0 B	13/07/2020, 18:04:22	1	128 MB	_SUCCESS
-rw-r--r--	alejandro	supergroup	0 B	13/07/2020, 18:04:21	1	128 MB	part-00000-5e3032c0-95cd-4dcb-af64-f46bc5a3b213-c000.csv
-rw-r--r--	alejandro	supergroup	12 B	13/07/2020, 18:04:15	1	128 MB	part-00003-5e3032c0-95cd-4dcb-af64-f46bc5a3b213-c000.csv
-rw-r--r--	alejandro	supergroup	7 B	13/07/2020, 18:04:15	1	128 MB	part-00004-5e3032c0-95cd-4dcb-af64-f46bc5a3b213-c000.csv
-rw-r--r--	alejandro	supergroup	11 B	13/07/2020, 18:04:15	1	128 MB	part-00005-5e3032c0-95cd-4dcb-af64-f46bc5a3b213-c000.csv
-rw-r--r--	alejandro	supergroup	11 B	13/07/2020, 18:04:15	1	128 MB	part-00006-5e3032c0-95cd-4dcb-af64-f46bc5a3b213-c000.csv
-rw-r--r--	alejandro	supergroup	23 B	13/07/2020, 18:04:15	1	128 MB	part-00011-5e3032c0-95cd-4dcb-af64-f46bc5a3b213-c000.csv

Figure 3.48: HDFS: Spark SQL Result in my_result

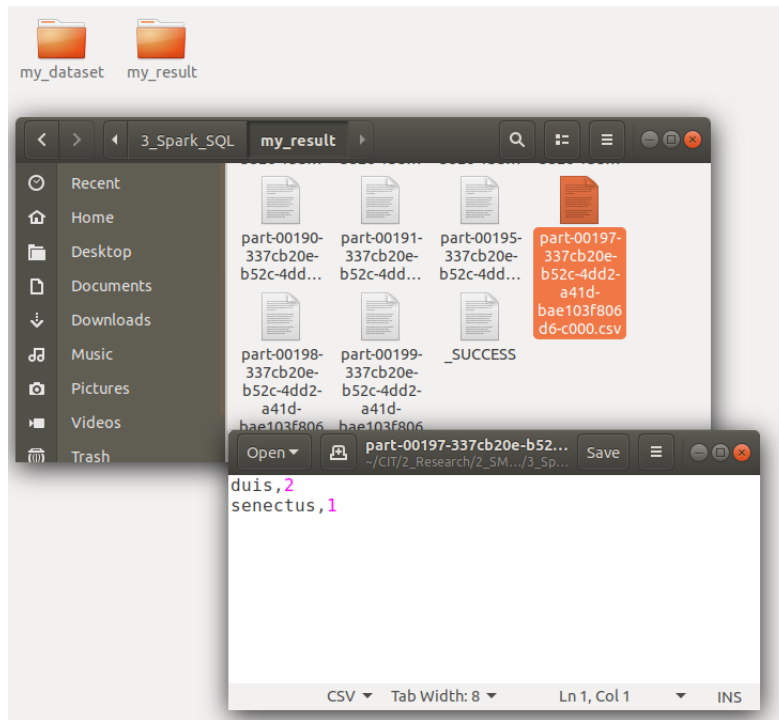


Figure 3.49: HDFS: Spark SQL Result Brought Back to Local File System

4 Conclusions and Future Work

In this deliverable we have introduced the Big Data ecosystem of tools we are going to use for storing and analysis bus and wind power-related large-datasets.

This ecosystem includes HDFS as a distributed file system (designed to efficiently allocate data across the multiple nodes of the cluster), Yarn as a resource manager (responsible for schedule and monitor the execution of our data analysis applications) and MapReduce, Spark Core and Spark SQL as frameworks for easily writing applications processing large-scale datasets across a cluster in a reliable, fault-tolerant manner. While Python is selected as the programming language of choice (with MapReduce, Spark Core and Spark SQL providing an API for it), the data analytics applications run on top of the Java Runtime Environment.

We have presented detailed scripts for installing, configuring and applying Java OpenJDK 8, Python 3.7.7, Hadoop 2.7.1 and Spark 2.4.5. In the case of Hadoop, the scripts include how to start and stop a Single Node Cluster with Pseudo-Distributed Operation. In the case of MapReduce, Spark Core and Spark SQL, a detailed explanation of an introductory example has been presented, together with a detailed explanation of its execution in the aforementioned cluster.

Bibliography

- [1] V. Mayer-Schonberger and K. Cukier, *Big Data*. Hmhbooks, 2013.
- [2] Apache Hadoop, <https://hadoop.apache.org/>.
- [3] K. Shvachko, H. Kuang, S. Radia, and R. Chansler, “The hadoop distributed file system,” in *26th Symposium on Mass Storage Systems and Technologies (MSST’10)*, 1–10. IEEE Computer Society, 2010.
- [4] Hadoop Distributed File System 3.2.1, <https://hadoop.apache.org/docs/r3.2.1/hadoop-project-dist/hadoop-hdfs/HdfsDesign.html>.
- [5] V. Vavilapalli, A. Murthy, C. Douglas, S. Agarwal, M. Konar, R. Evans, T. Graves, J. Lowe, H. Shah, S. Seth, B. Saha, C. A. Curino, O. O’Malley, S. R. Radia, B. C. Reed, and E. Balde-schieler, “Apache hadoop yarn: yet another resource negotiator,” in *4th Symposium on Cloud Computing (SOCC’13)*, 1–16. ACM, 2013.
- [6] Apache Hadoop Yarn, <https://hadoop.apache.org/docs/current/hadoop-yarn/hadoop-yarn-site/YARN.html/>.
- [7] J. Dean and S. Ghemawat, “Mapreduce: Simplified data processing on large clusters,” in *6th Symposium on Operating System Design and Implementation (OSDI’04)*, 137–150, 2004.
- [8] Hadoop MapReduce 3.2.1, <https://hadoop.apache.org/docs/r3.2.1/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html>.
- [9] Apache Spark, <https://spark.apache.org/>.
- [10] M. Zaharia, M. Chowdhury, T. Das, A. Dave, J. Ma, M. McCauly, M. J. Franklin, S. Shenker, and I. Stoica, “Resilient distributed datasets: A fault-tolerant abstraction for in-memory cluster computing,” in *9th Symposium on Networked Systems Design and Implementation (NSDI’12)*, 15–28. USENIX, 2012.
- [11] Apache Spark Core, <https://spark.apache.org/docs/2.4.5/rdd-programming-guide.html>.
- [12] M. Armbrust, R. S. Xin, C. Lian, Y. Huai, D. Liu, J. K. Bradley, X. Meng, T. Kaftan, M. J. Franklin, A. Ghodsi, and M. Zaharia, “Spark sql: Relational data processing in spark,” in *15th International Conference on Management of Data (SIGMOD’15)*, 1383–1394. ACM Press, 2015.
- [13] Apache Spark SQL, <https://spark.apache.org/docs/2.4.5/sql-programming-guide.html>.
- [14] Java OpenJDK, <https://openjdk.java.net/>.
- [15] Hadoop Streaming 2.7.1, <https://hadoop.apache.org/docs/r2.7.1/hadoop-streaming/HadoopStreaming.html>.
- [16] The Scala Programming Language, <https://www.scala-lang.org/>.
- [17] Ubuntu 20.04 LTS, <https://ubuntu.com/>.
- [18] Java OpenJDK 8, <https://openjdk.java.net/projects/jdk8u/>.
- [19] Java OpenJDK 11, <https://openjdk.java.net/projects/jdk11/>.
- [20] Python 3.7.7, <https://www.python.org/downloads/release/python-377/>.
- [21] Python 3.8.2, <https://www.python.org/downloads/release/python-382/>.

- [22] pip 20.1, <https://pypi.org/project/pip/>.
- [23] Hadoop 2.7.1, <https://hadoop.apache.org/docs/r2.7.1/>.
- [24] Hadoop 3.2.1, <https://hadoop.apache.org/docs/r3.2.1/>.
- [25] OpenSSH, <https://www.openssh.com/>.
- [26] Lorem Ipsum Generator, <https://www.lipsum.com/>.
- [27] Apache Spark 2.4.5, <https://spark.apache.org/docs/2.4.5/>.
- [28] Apache Spark 3.0.0, <https://spark.apache.org/docs/3.0.0-preview/>.
- [29] pyspark 2.4.5, <https://pypi.org/project/pyspark/>.
- [30] PyCharm 2020.1, <https://www.jetbrains.com/pycharm/>.



SMART electric Buses

July 15, 2020